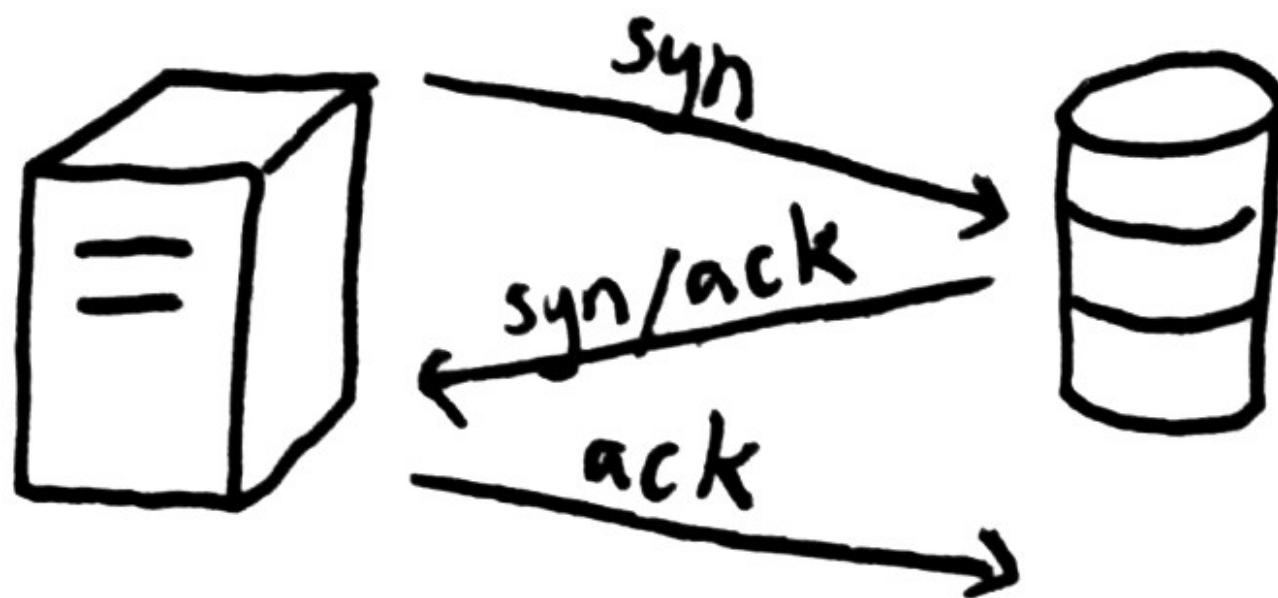


HEY,

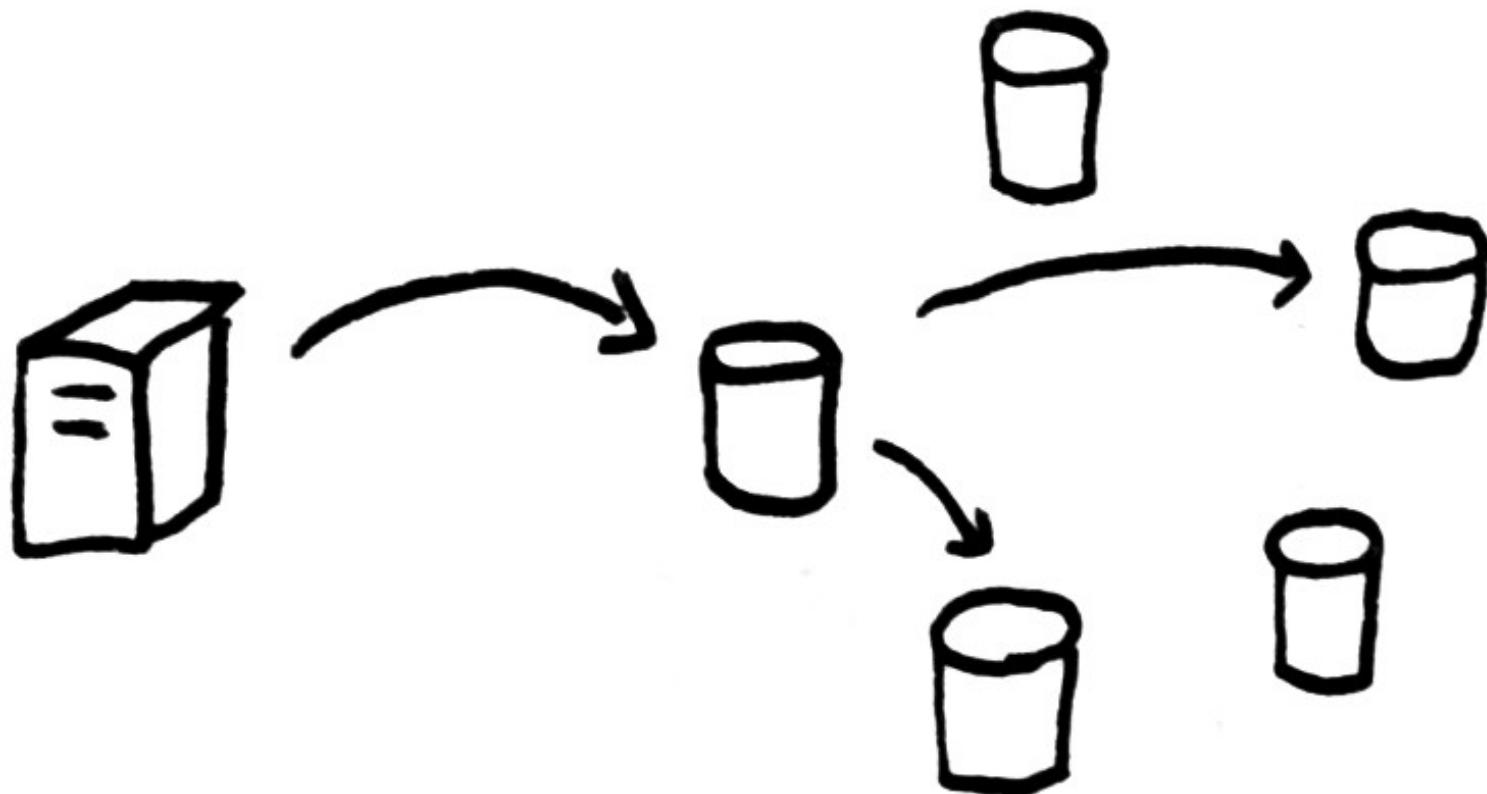
I just met you



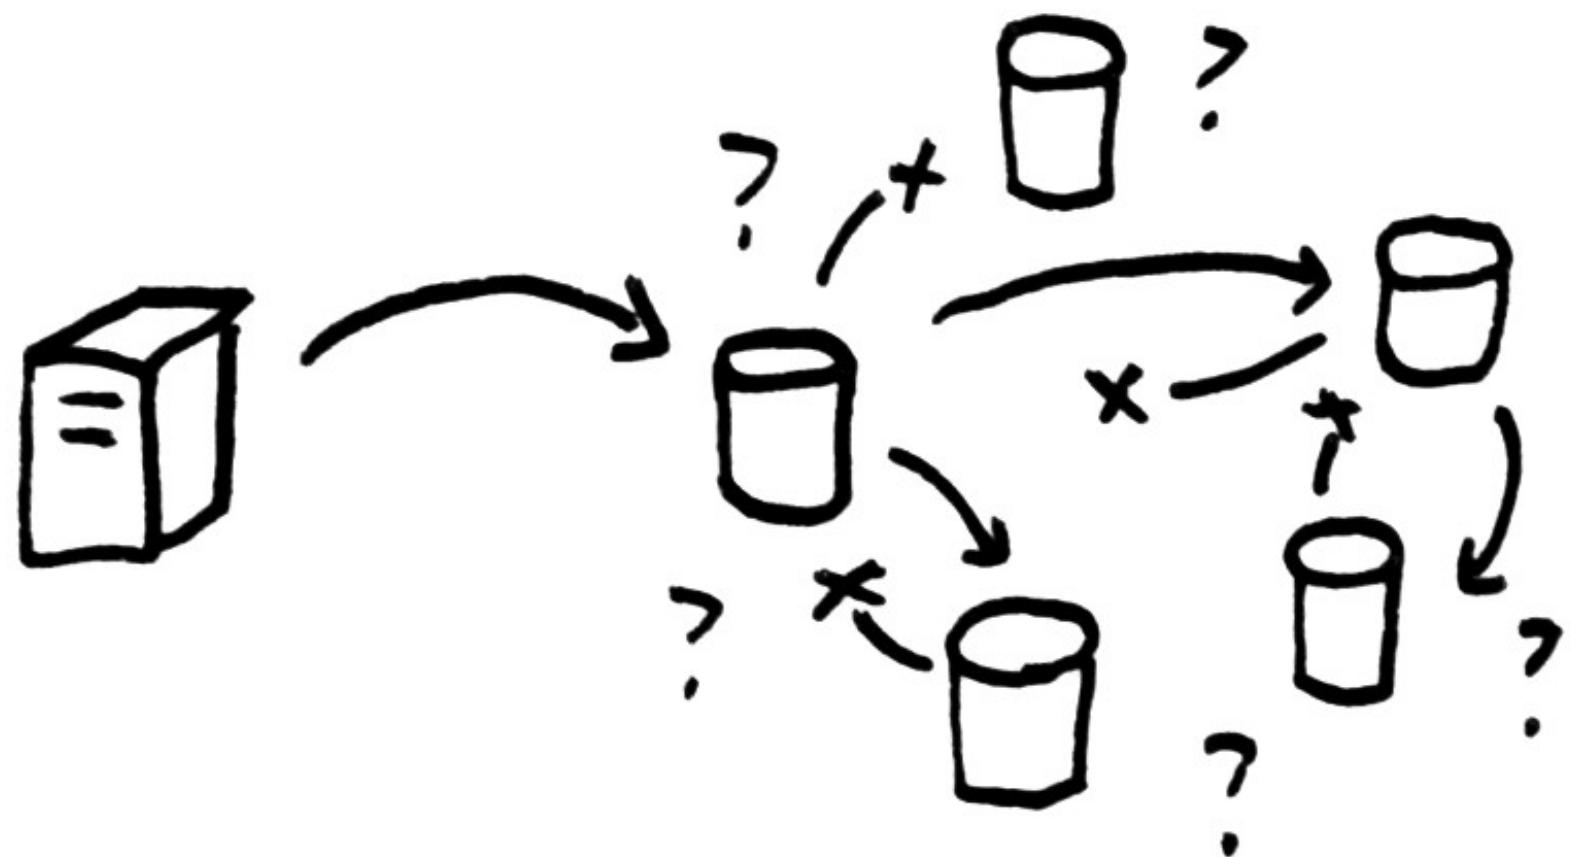
Our network's flaky



But here's my data



So store it maybe?



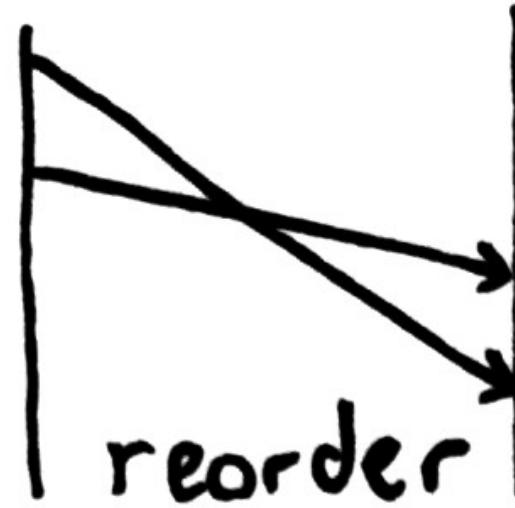
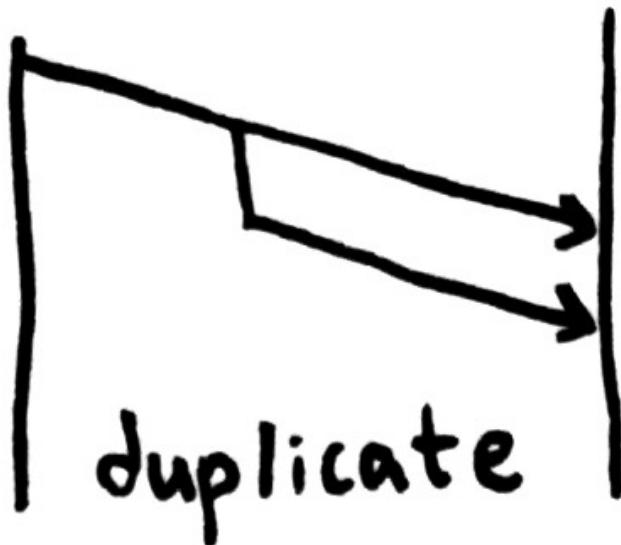
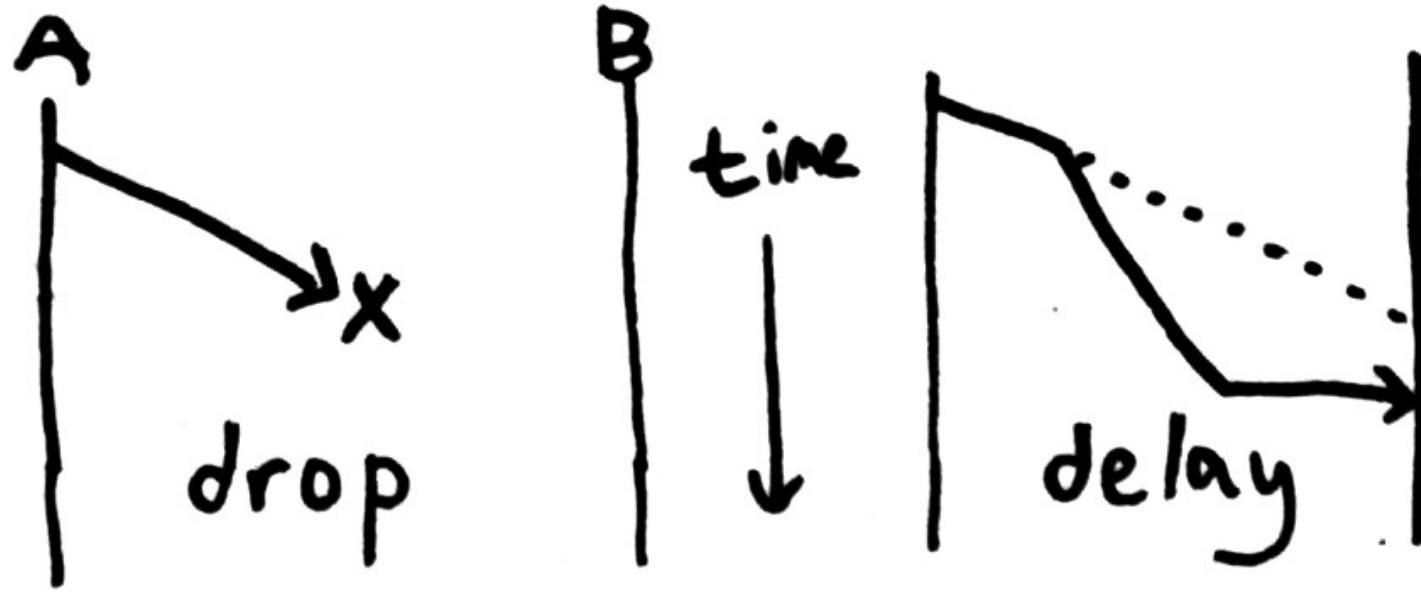


@aphyr
(Kyle Kingsbury)

PARTITIONS
for
EVERYONE!

What are
partitions,
anyway?

Formally:



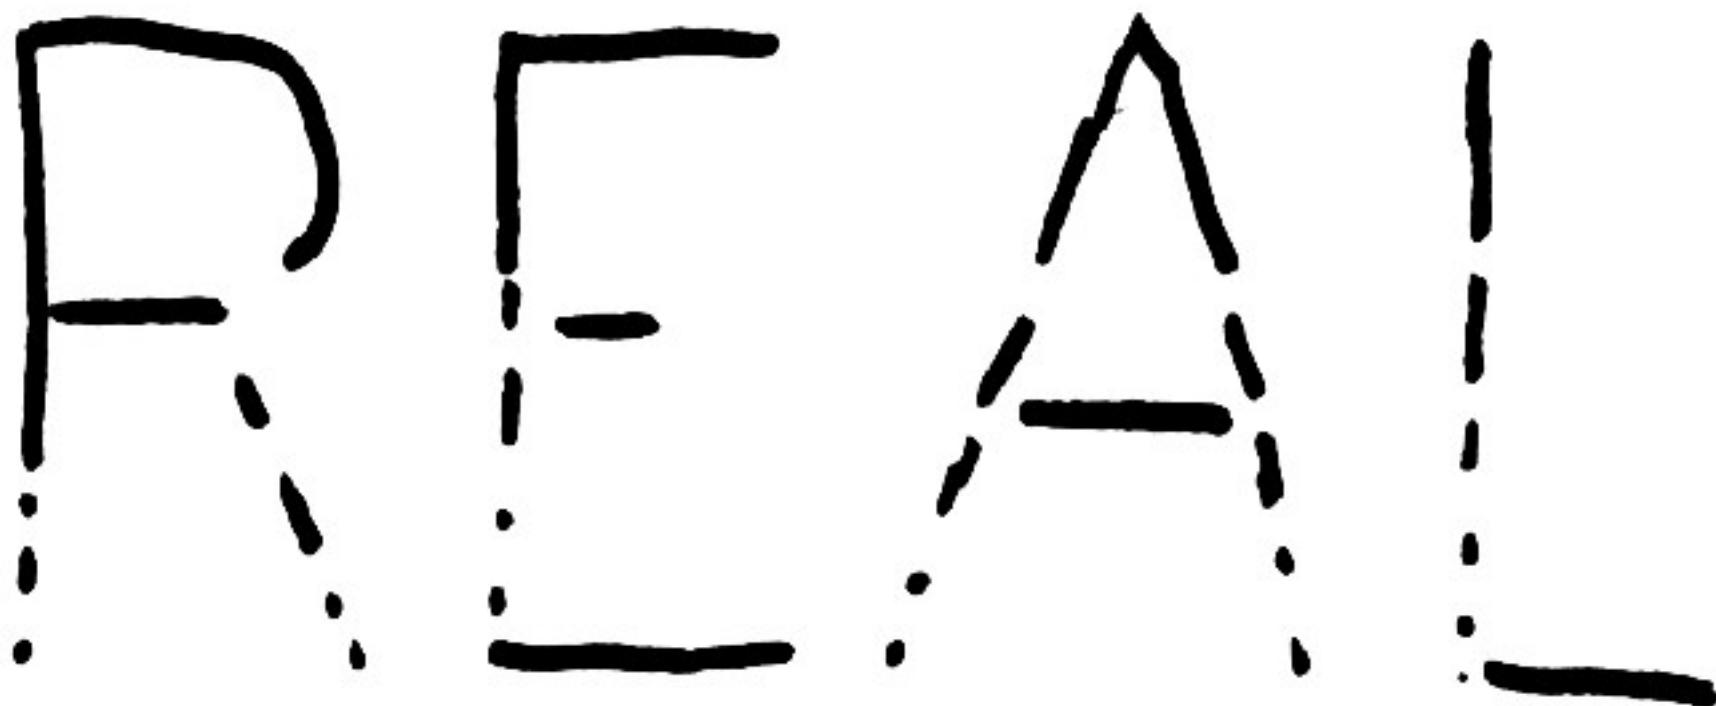
IP networks do all four.

TCP means no dupes, reorders

- Unless you retry!

Delays are indistinguishable
from drops, after timeout

Partitions are



Microsoft : 40.8 link failures/day

- 5 minutes — one week

- Redundancy improves packet loss
by only 43%.

Google: in 1 cluster...

- 5 racks see 50% pkt loss
- 8 net maint/yr, 30 min loss in 4
- 3 router failures/yr

Yahoo Sherpa

Amazon Dynamo

Google Chubby

Causes

- GC pauses
- Network maint
- Segfaults & crashes
- Faulty NICs
- Bridge loops

... continued

- Spanning tree
- Hosted networks
- The cloud
- WAN links & Backbones

Cloudflare

Pagerduty

Netflix

Fog Creek

Github

City Cloud

Searchbox.io

AWS

Freistel IT

Twilio



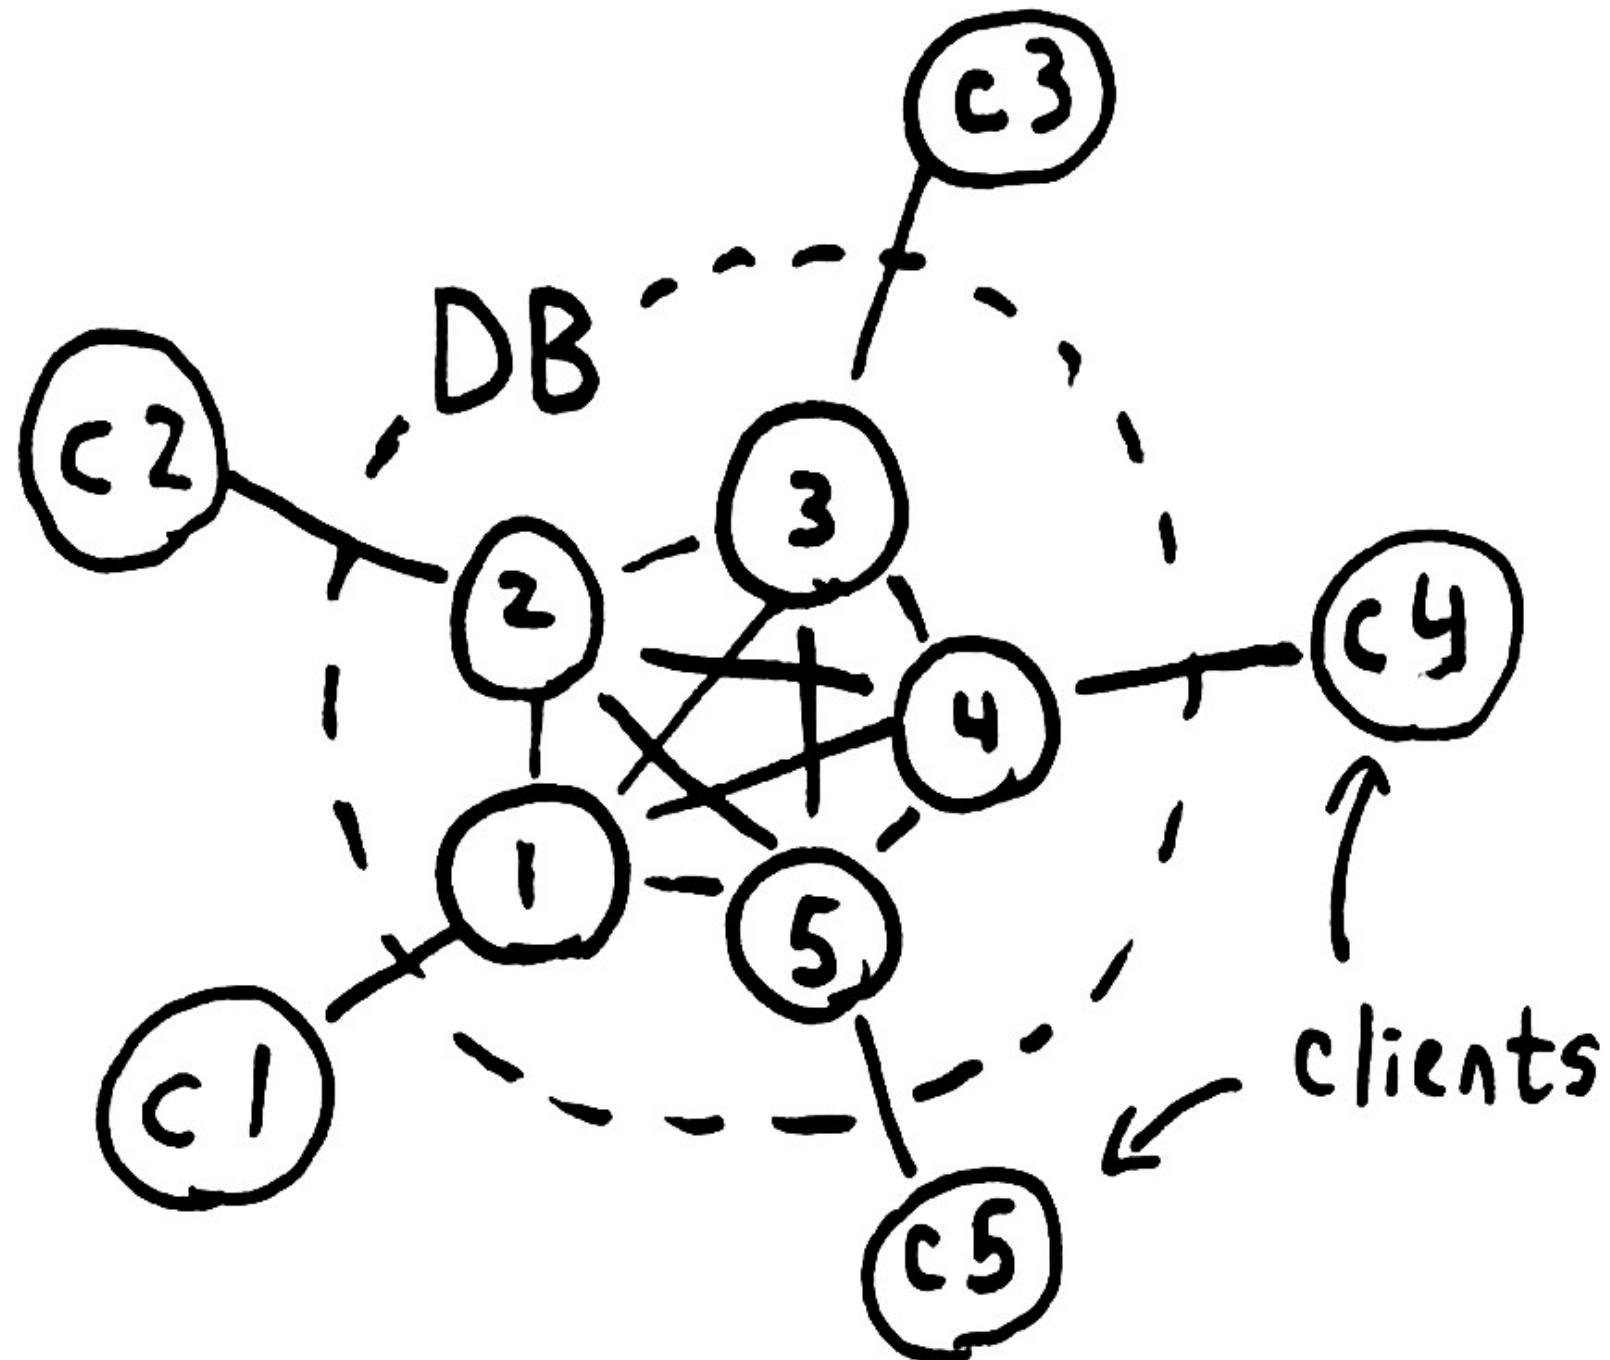
PLANET
TERAHI
(you are not here)

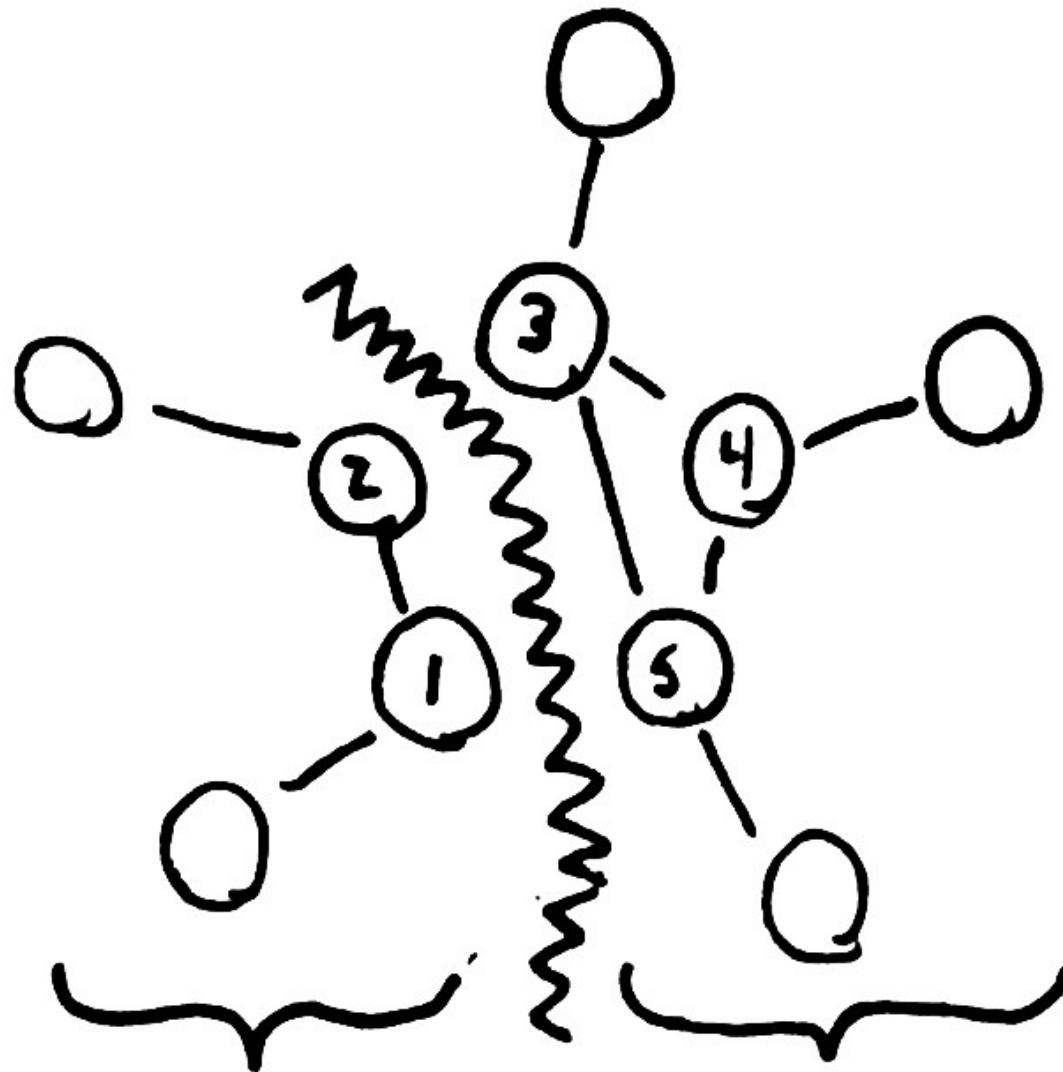
Distributed Systems

tend to fail

part... partially

Jepsen is an
exploration of
those failure modes

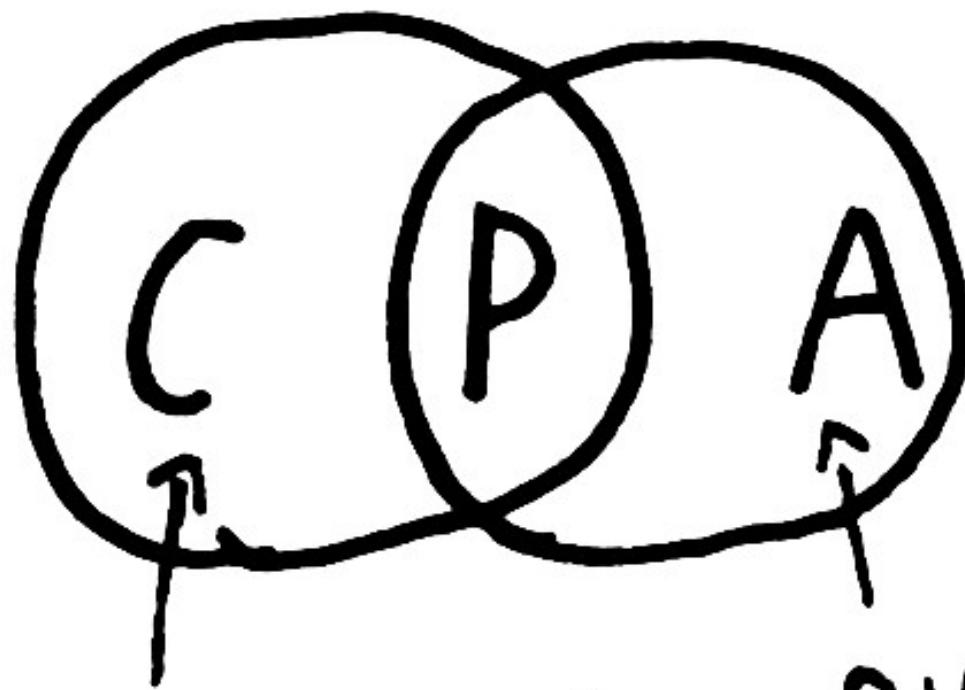




component component

(minority) (majority)

In broad terms



serializability availability

Both AP and CP
systems can remain
available in this case

CP

needs a majority
or primary node -
forces re-election

AP

doesn't care about
partitions – business
as normal

- # Single - server Postgres
- CP
 - SPOF
 - consistent
 - clients may not agree

Redis (w/failover)

- Split-brain
- Drops all data on one component
- Not even close to consistent.

MongoDB

- All consistency levels led to lost data
- Weak Defaults
- Majority was a bug, fixed

Riak

- LWW means dropping writes
- Even with PR/PW = ALL
- Even with lock service
- use CRDTs

That was Then,

This is  meow

That was Then,

This is  now

Zookeeper

NuoDB

Kafka

Cassandra

ZIK

- Paxos - alike protocol
- Linearizable writes
- Majority w/ leaders

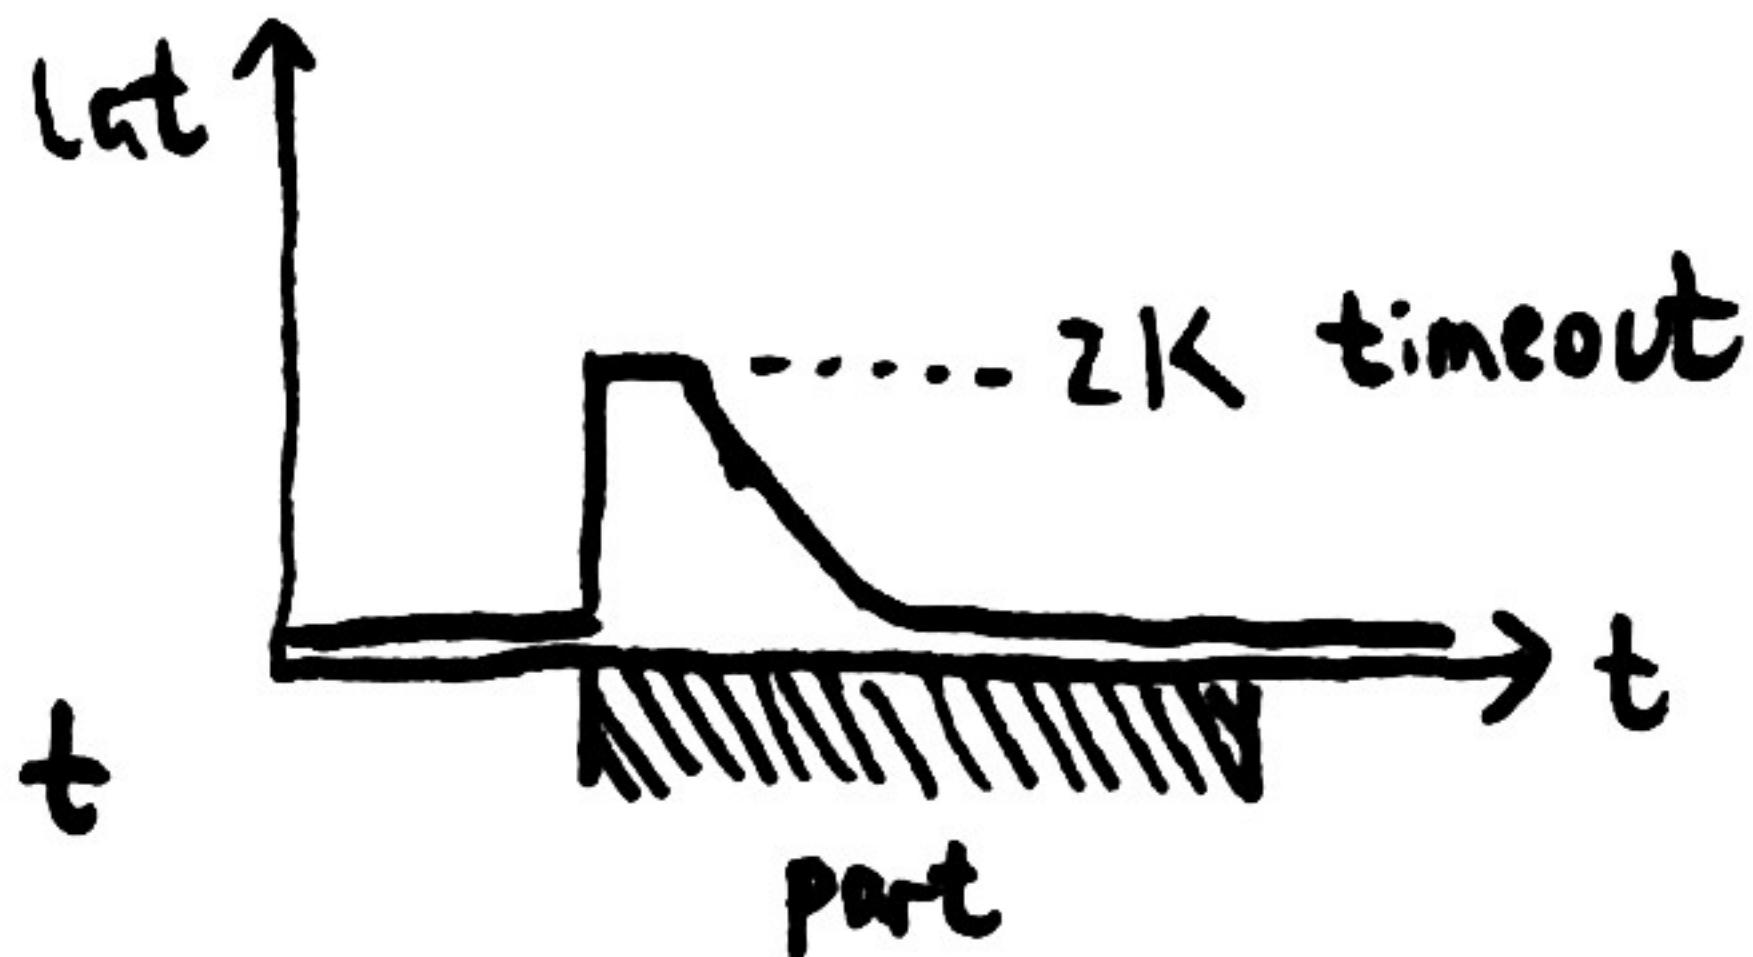
- Detects partition quickly
- Writes immediately fail
- Majority component continues after election
- Only partly available

Boooring.

```
121 :error 9196
122 :ok 13211
123 :ok 14312
124 :ok 15154
125 :error 8693
126 :error 8696
127 :ok 12884
128 :ok 14014
129 :ok 13922
130 :error 8194
131 :error 8195
132 :ok 12280
133 :ok 13237
134 :ok 13898
135 :error 7695
136 :error 7697
137 :ok 11408
138 :ok 12836
139 :ok 13373
140 :error 11576
141 :error 10634
142 :ok 10858
143 :ok 11785
144 :ok 10070
145 :error 11076
146 :error 10135
147 :ok 10875
148 :ok 11401
149 :ok 12642
150 :error 10576
151 :error 11213
152 :ok 10074
153 :ok 11725
154 :ok 11850
155 :error 11454
156 :error 10820
157 :ok 8944
158 :ok 10326
159 :ok 10071
160 :error 9576
161 :error 10213
162 :ok 8243
163 :ok 9684
164 :ok 11150
165 :error 9076
166 :error 8133
167 :ok 8220
168 :ok 10041
169 :ok 10644
170 :error 9955
171 :error 7635
172 :ok 8121
```

```
49465 [pool-11-thread-20-SendThread(n1:2181)] INFO org.apache.zookeeper.ClientCnxn - Socket connection established to n1/10.10.3.242:2181, initiating session
49469 [pool-11-thread-20-SendThread(n1:2181)] INFO org.apache.zookeeper.ClientCnxn - Unable to read additional data from server sessionid 0x0, likely server has closed socket, closing socket connection and attempting reconnect
49548 [pool-13-thread-39] INFO org.apache.zookeeper.ZooKeeper - Session: 0x0 closed
49548 [pool-13-thread-36-EventThread] INFO org.apache.zookeeper.ClientCnxn - EventThread shut down
49548 [pool-13-thread-39] INFO org.apache.zookeeper.ZooKeeper - Initiating client connection, connectString=n2:2181 sessionTimeout=2000 watcher=org.apache.curator.ConnectionState@110d66fa
49549 [pool-13-thread-39-SendThread(n2:2181)] INFO org.apache.zookeeper.ClientCnxn - Opening socket connection to server n2/10.10.3.52:2181. Will not attempt to authenticate using SASL (unknown error)
49550 [pool-13-thread-39-SendThread(n2:2181)] INFO org.apache.zookeeper.ClientCnxn - Socket connection established to n2/10.10.3.52:2181, initiating session
49550 [pool-13-thread-39-SendThread(n2:2181)] INFO org.apache.zookeeper.ClientCnxn - Unable to read additional data from server sessionid 0x0, likely server has closed socket, closing socket connection and attempting reconnect
300 :error 11696
301 :error 9640
302 :ok 66
303 :ok 67
304 :ok 50
305 :error 11196
306 :error 9140
307 :ok 44
308 :ok 63
309 :ok 50
310 :error 10696
311 :error 9970
312 :ok 17
313 :ok 46
314 :ok 13
315 :error 10195
50777 [pool-13-thread-39-SendThread(n2:2181)] INFO org.apache.zookeeper.ClientCnxn - Opening socket connection to server n2/10.10.3.52:2181. Will not attempt to authenticate using SASL (unknown error)
50778 [pool-13-thread-39-SendThread(n2:2181)] INFO org.apache.zookeeper.ClientCnxn - Socket connection established to n2/10.10.3.52:2181, initiating session
50788 [pool-13-thread-39-SendThread(n2:2181)] INFO org.apache.zookeeper.ClientCnxn - Unable to read additional data from server sessionid 0x0, likely server has closed socket, closing socket connection and attempting reconnect
316 :error 11514
317 :ok 50
318 :ok 50
319 :ok 77
320 :error 9586
321 :error 8969
322 :ok 45
323 :ok 75
324 :ok 51
51374 [pool-11-thread-42] WARN org.apache.curator.ConnectionState - Connection attempt unsuccessful after 2026 (greater than max timeout of 2000). Resetting connection and trying again with a new connection.
51397 [pool-11-thread-20-SendThread(n1:2181)] INFO org.apache.zookeeper.ClientCnxn - Opening socket connection to server n1/10.10.3.242:2181. Will not attempt to authenticate using SASL (unknown error)
51397 [pool-11-thread-20-SendThread(n1:2181)] INFO org.apache.zookeeper.ClientCnxn - Socket connection established to n1/10.10.3.242:2181, initiating session
51412 [pool-13-thread-22] WARN org.apache.curator.ConnectionState - Connection attempt unsuccessful after 2037 (greater than max timeout of 2000).
```

Latency



Zookeeper

78% acknowledged

0% false positives

0.1% false negatives

ZK strategies

- Use it,
- Use curator, etc.
- Consider linearizability
carefully! (writes only)

NUODB

Cloud Computing ›

Brewer's CAP Conjecture is False

57 posts by 14 authors  

and a purported "proof" at

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.20.1495&rep=rep1&type=pdf>

The CAP conjecture, I am convinced, is false and can be proven false.

The CAP conjecture has been a theoretical millstone around the neck of all ACID systems. Good riddance.

This is the first wooden stake for the heart of the noSQL movement.
There are more coming.

--

Jim Starkey

Founder, NimbusDB, Inc.

978 526-1376

If the CAP theorem means that all surviving nodes must be able to continue processing without communication after a network failure, then NUODB is not partition resistant.

If partition resistance includes the possibility for a surviving subset of the chorus to sing on, then NUODB refutes the CAP theorem.

- Design is complex
- Monotonic leaders for ensembles over each row, table, db, ...
- Claims to be C, A, 3P?

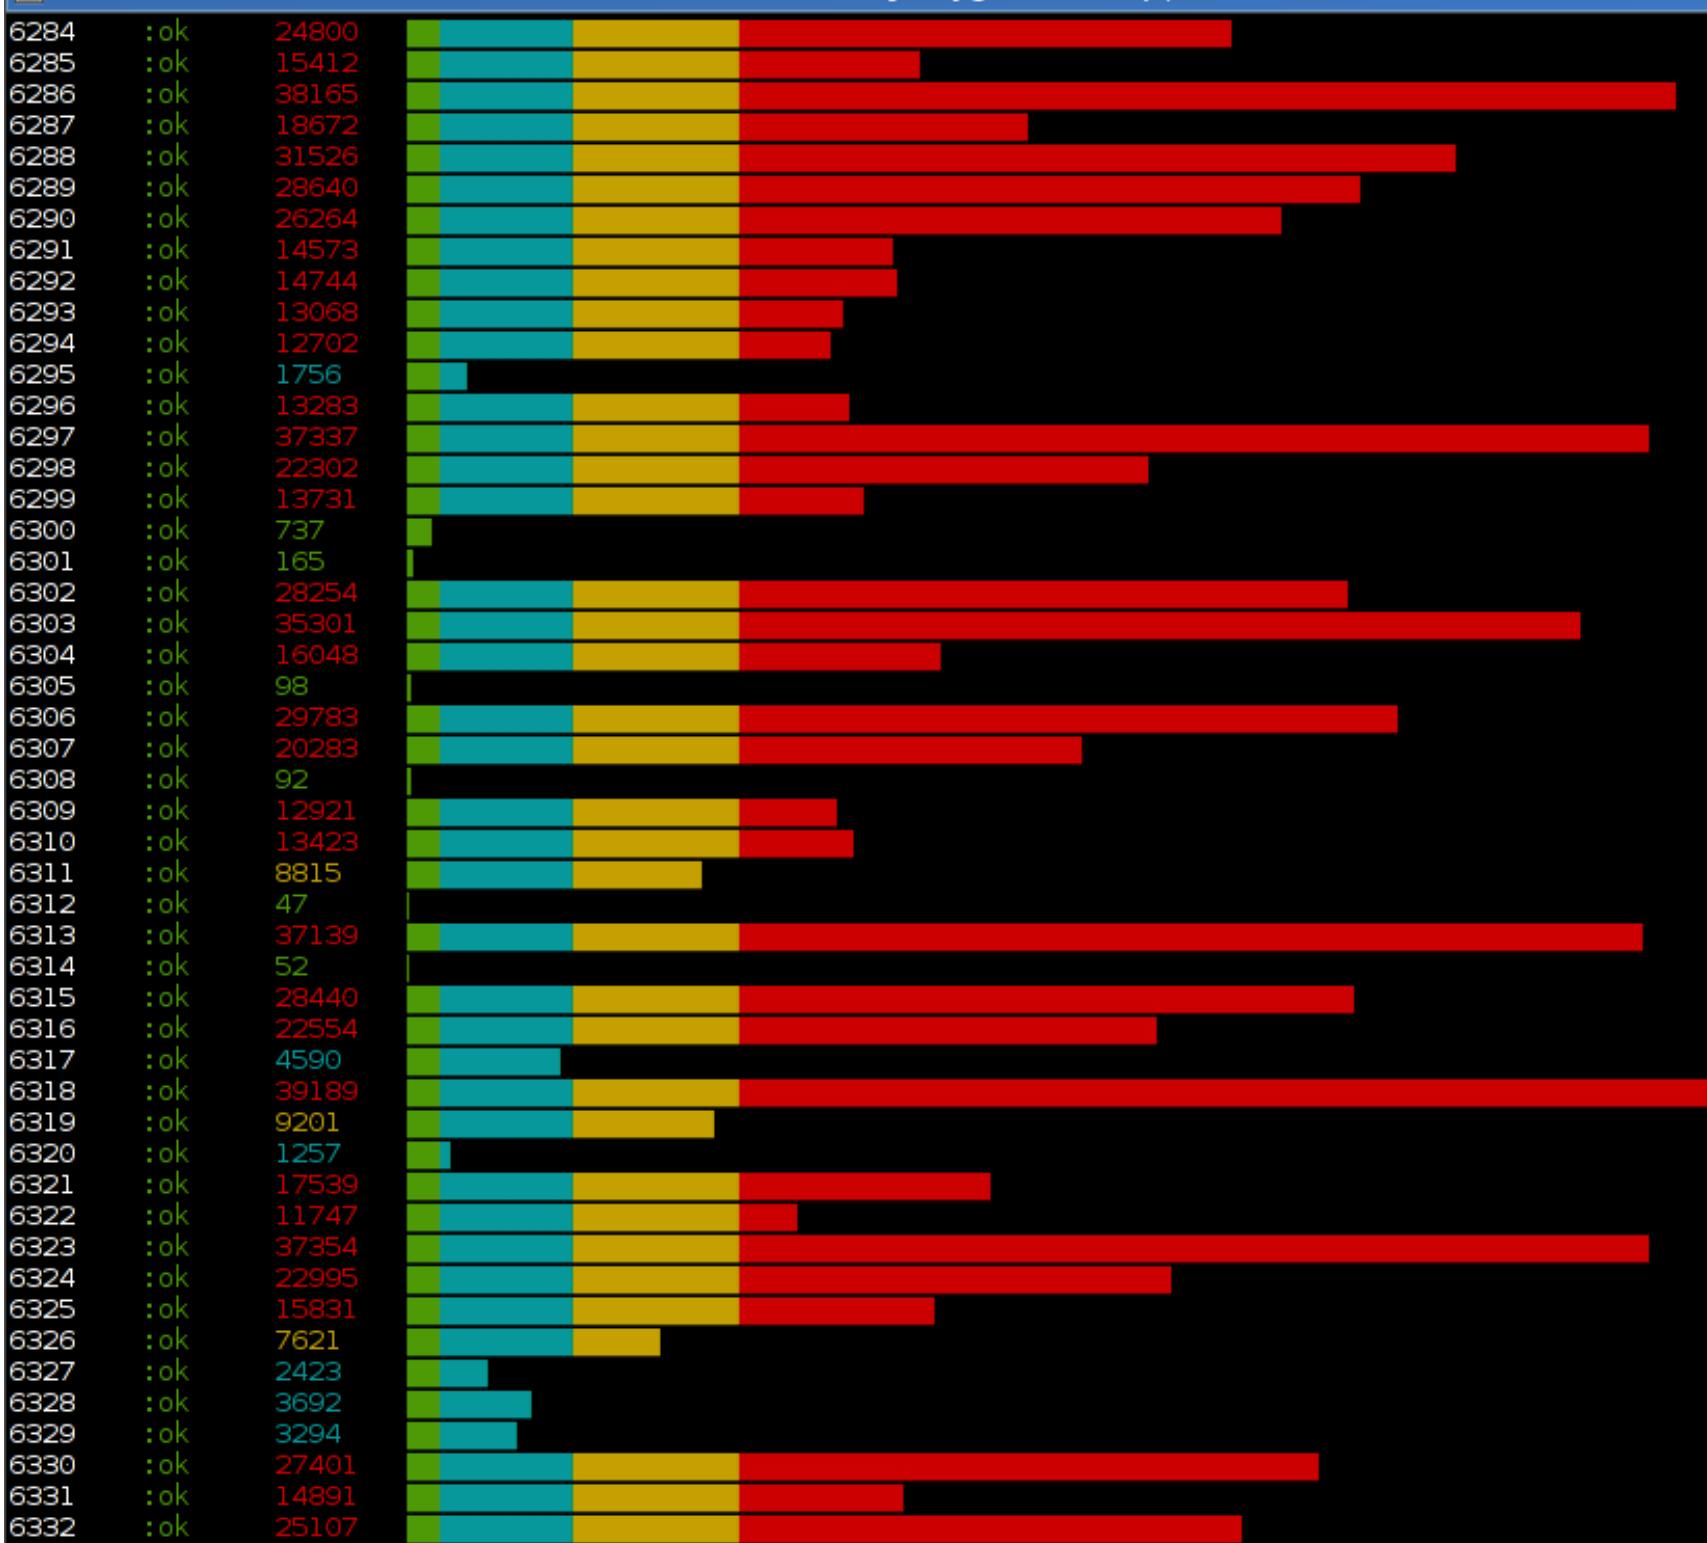
```
SQL> DROP TABLE SET_APP CASCADE;
sequence "SYSTEM.SET_APP$IDENTITY_SEQUENCE" does not exist
SQL> CREATE SEQUENCE SYSTEM.SET_APP$IDENTITY_SEQUENCE;
sequence SYSTEM.SET_APP$IDENTITY_SEQUENCE already exists
SQL> █
```

- Race conditions in join
- Can self-join nodes
- Race conditions in DB create
- Basically impossible to start reliably
- Does not bring crashed node storage back online

"Duplicate value in unique
index SEQUENCES..PRIMARY_KEY

- End of stream reached
- Broken pipe,
- Conn reset
- indefinite latency

```
2483 :ok    44
2484 :ok    43
2485 :ok    51
2486 :ok    30
2487 :ok    69
2488 :ok    39
2489 :ok    86
2490 :ok    29
2491 :ok    68
2492 :ok    24
2493 :ok    44
2494 :ok    45
2495 :ok    37
2496 :ok    98
2497 :ok    31
2498 :ok    54
2499 :ok    62
2500 :ok    83
2501 :ok    49
2502 :ok   112
2503 :ok    64
2504 :ok    88
2505 :error 65000
2506 :error 65000
2507 :error 65000
2508 :error 65000
2509 :error 65001
2510 :error 65000
2511 :error 65000
2512 :error 65000
2513 :error 65000
2514 :error 65000
2515 :error 65011
2516 :error 65000
2517 :error 65000
2518 :error 65000
2519 :error 65011
2520 :error 65000
2521 :error 65000
2522 :error 65000
2523 :error 65000
2524 :error 65003
2525 :error 65003
2526 :error 65000
2527 :error 65000
2528 :error 65000
2529 :error 65009
2530 :error 65000
2531 :error 65000
```



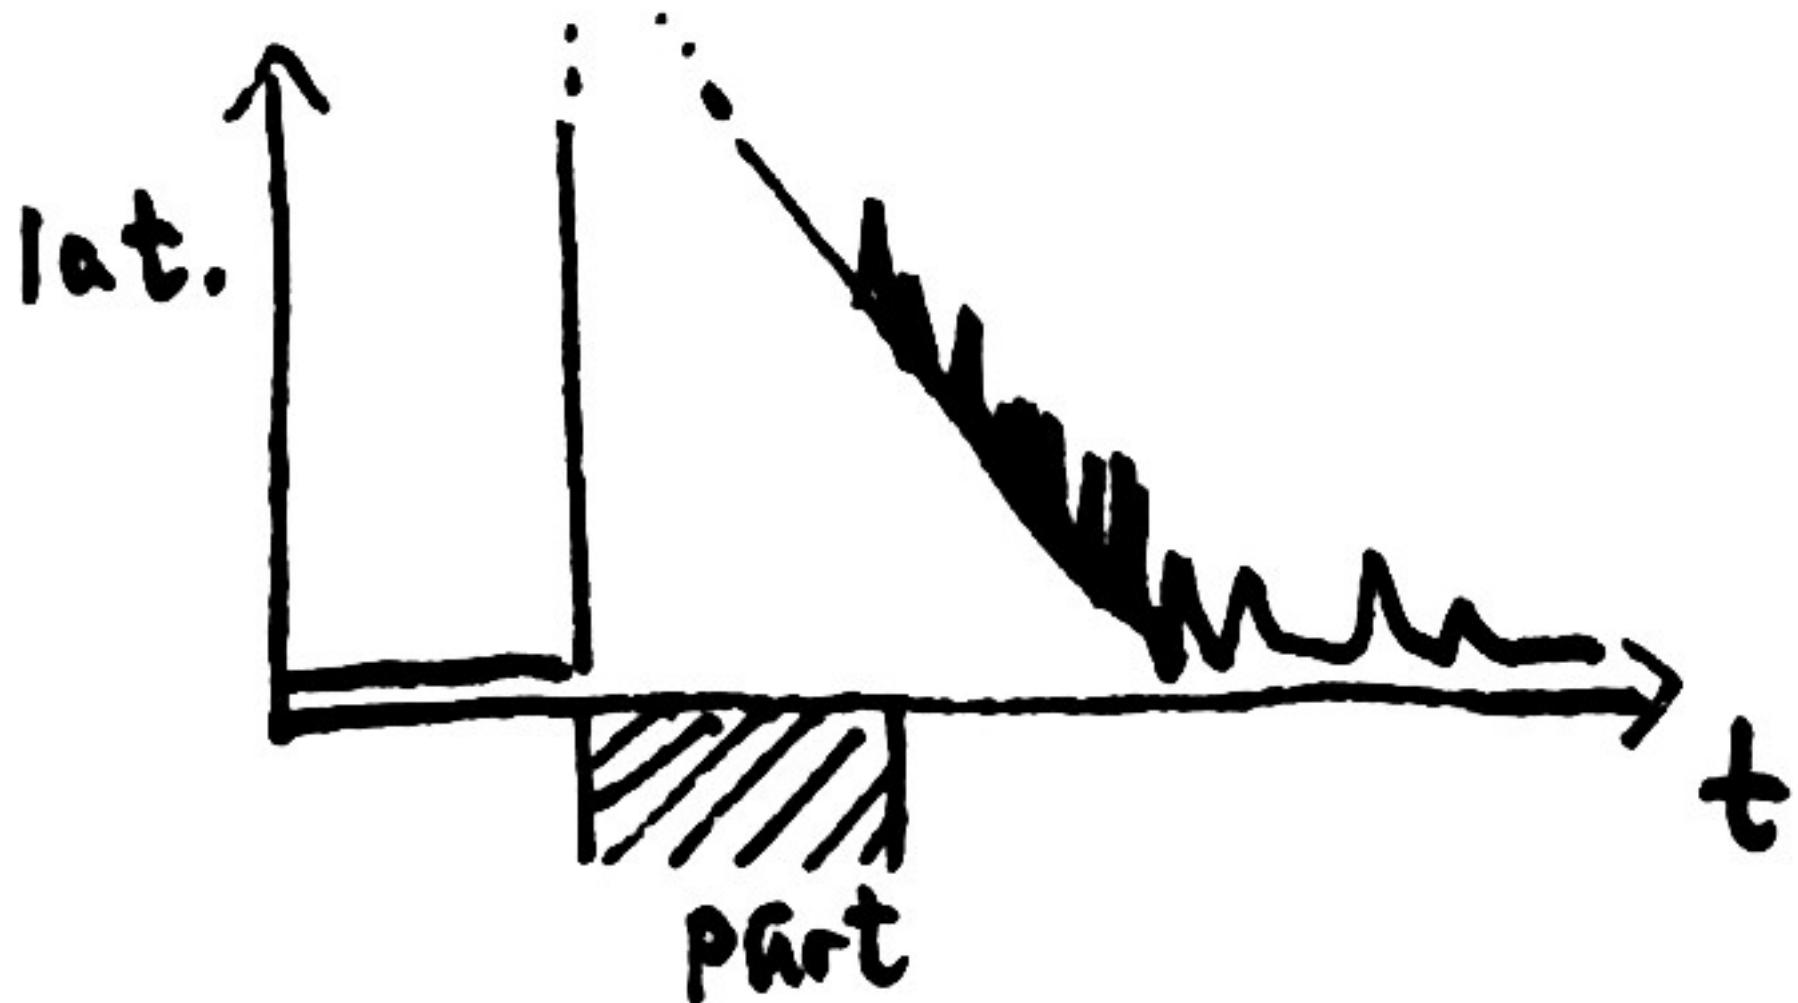
NuoDB

55% acknowledged

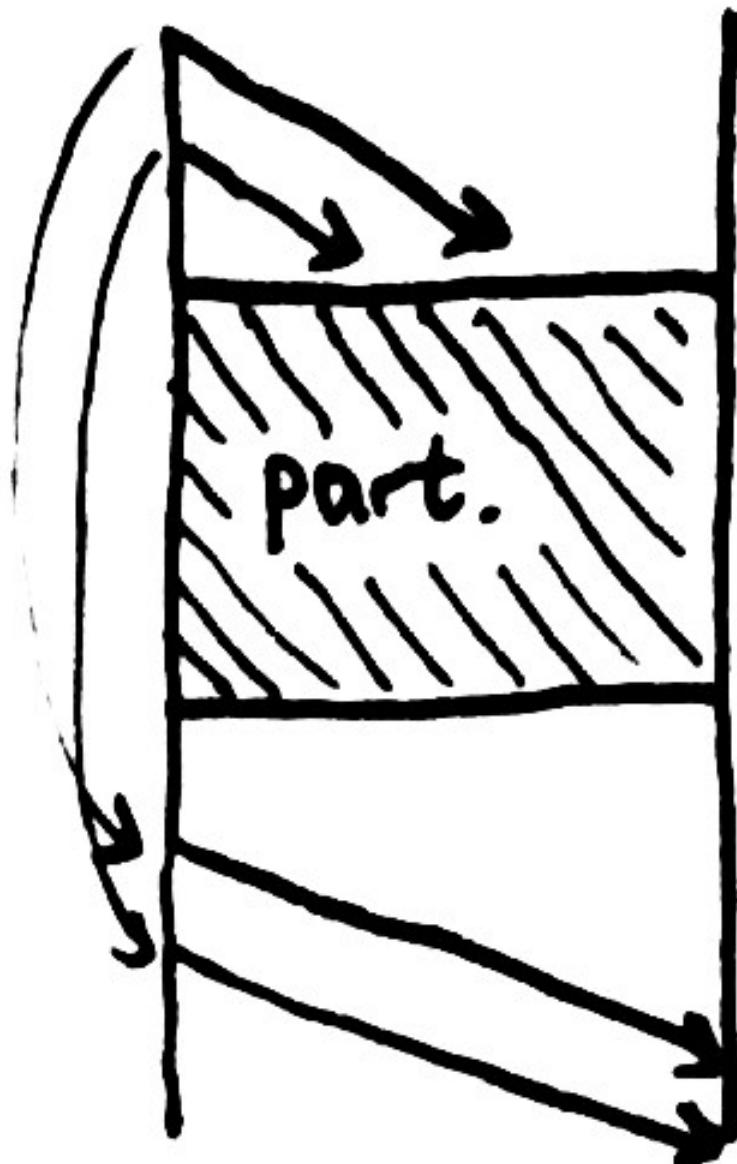
0% false positives

0.1% false negatives

Latency



"Is a database
which takes five
minutes to respond
really available?"



This is just
pushing the
problem off
for later.

NuoDB is working on
LIVENESS DETECTION

and many of these bugs.

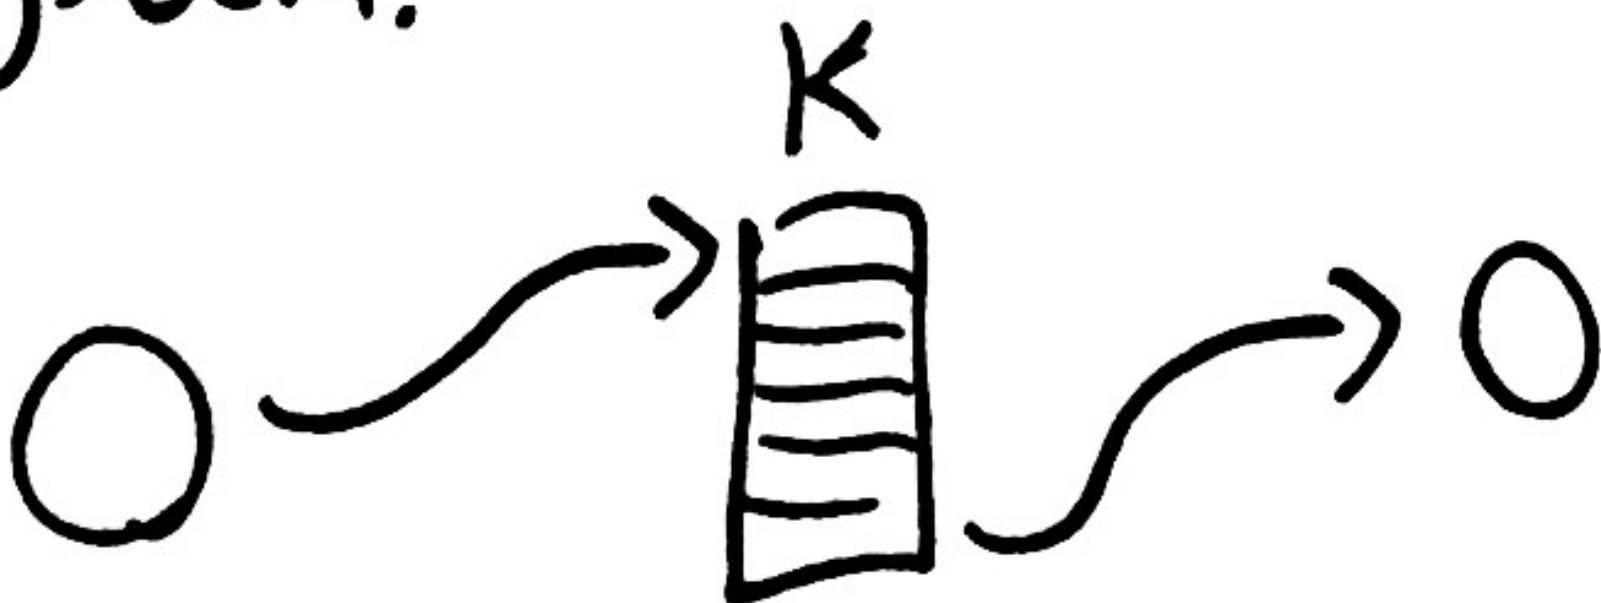
Stay tuned!

NuoDB strategies

- It's still young
- Enable liveness detection
- Wait 'n' see

KAFKA

Kafka is a messaging system!



Provides a sharded
immutable log of
messages.

New in 0.8.0:

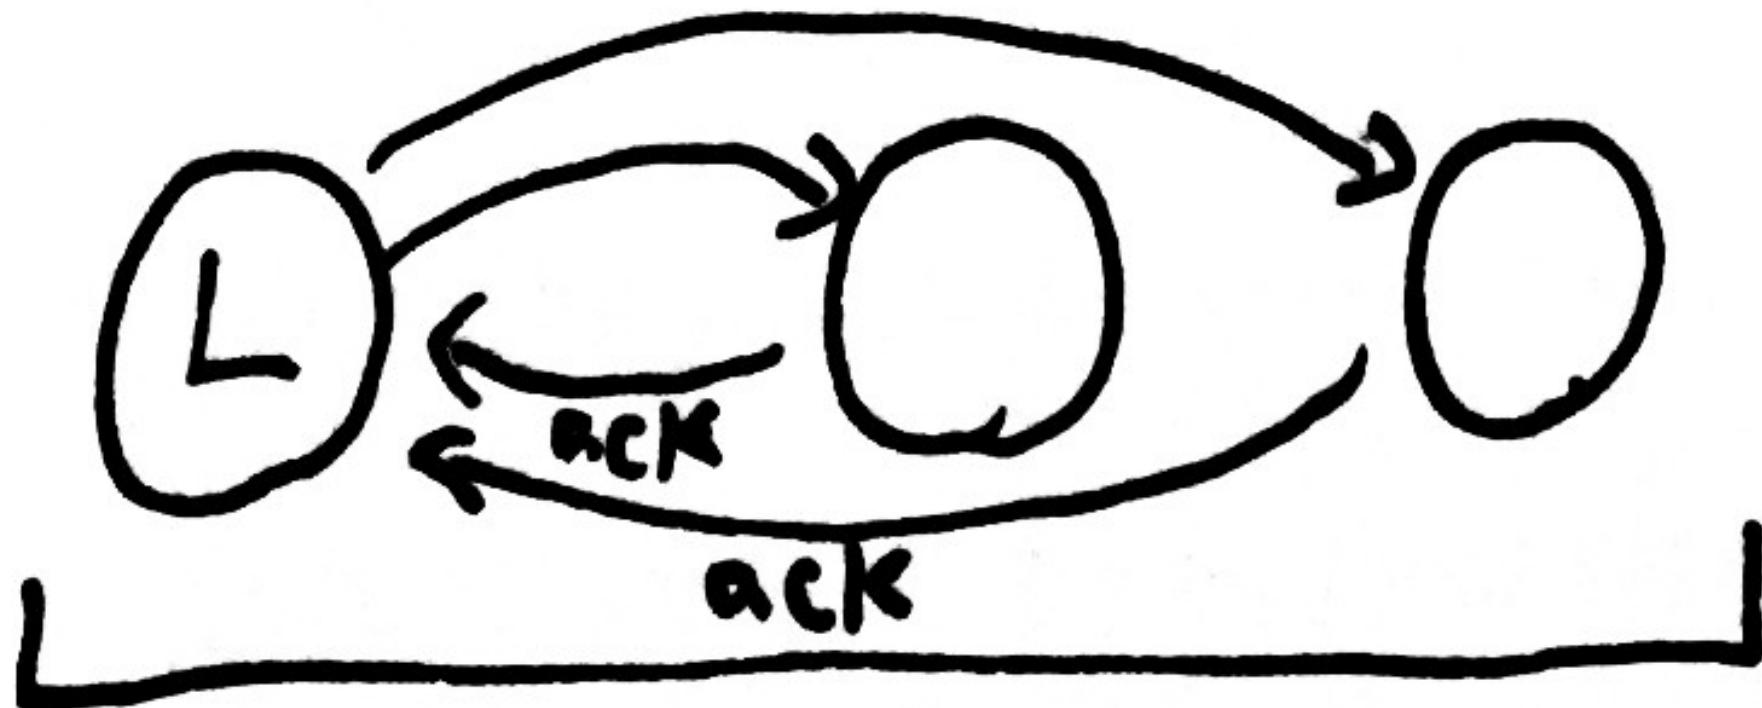
REPLICATION!

[!\[\]\(c306a97f2b1bda03752a7a3f55013029_img.jpg\) Email](#) [!\[\]\(c9dfecd68d2224a37fe045ee90b52ee5_img.jpg\) Like](#) [!\[\]\(2d10ff32d96230898a87476112586b16_img.jpg\) Save](#) [!\[\]\(0bfab6a86a7b1d12ac32e93bc1f78a2e_img.jpg\) Embed](#)

Kafka Replication: Pick CA

- Brokers within a datacenter
 - i.e., network partitioning is rare
- Strong consistency
 - replicas byte-wise identical
- Highly available
 - typical failover time: < 10ms

In a shard:



"in-sync replicas"



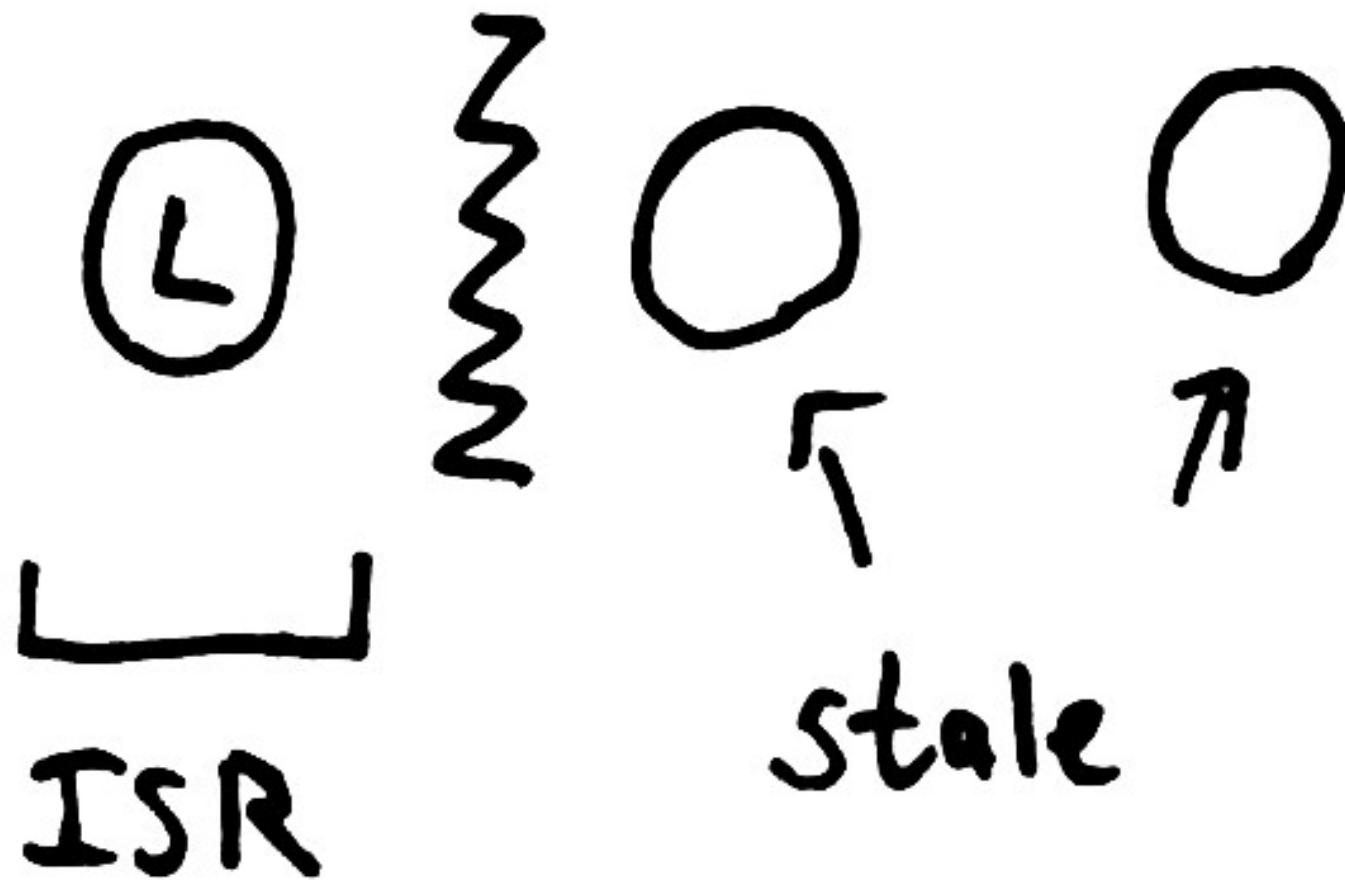
" f replicas means:

\Rightarrow tolerate

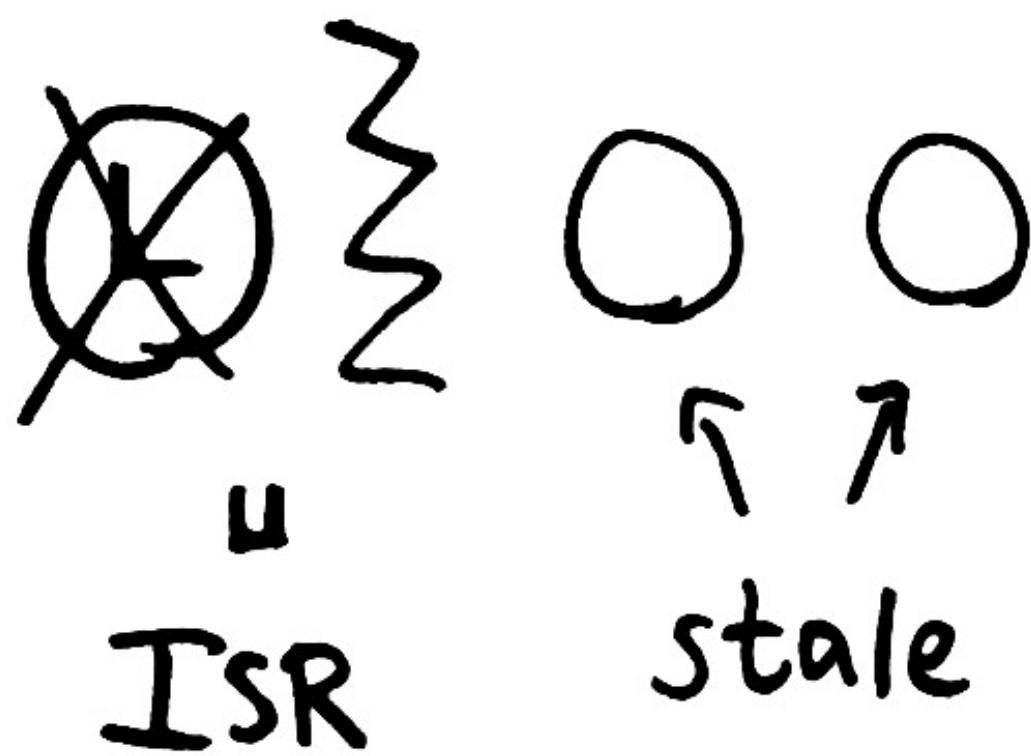
$f - 1$

"failures"

How?



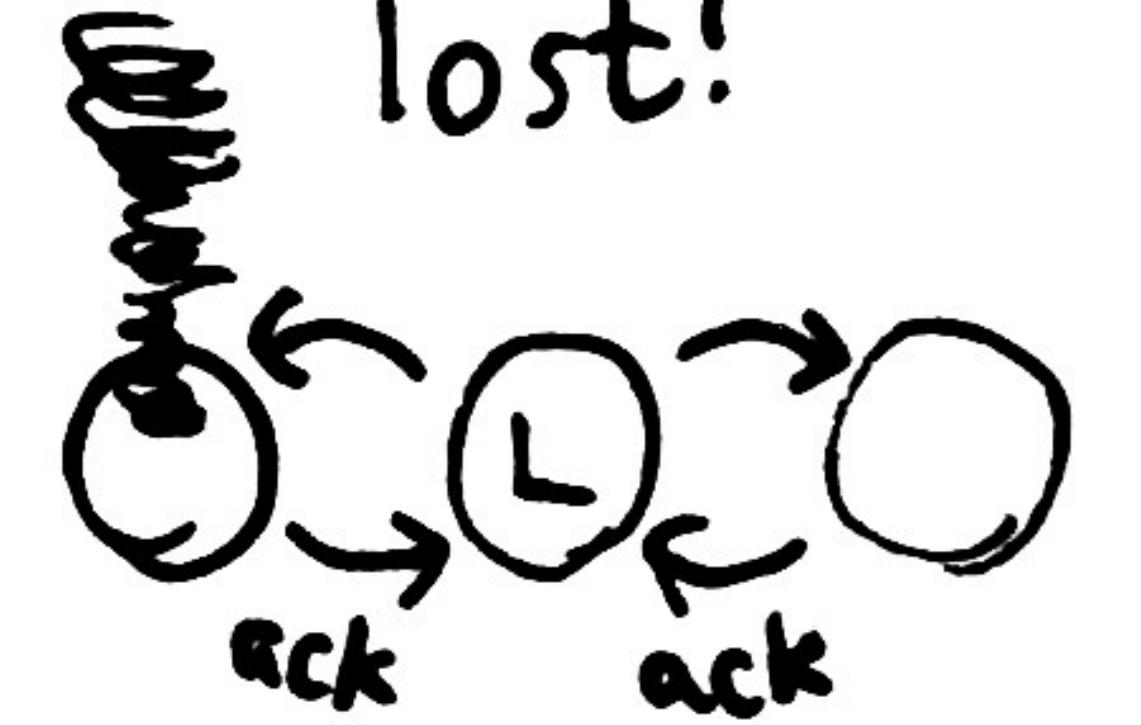
What if leader disappears?



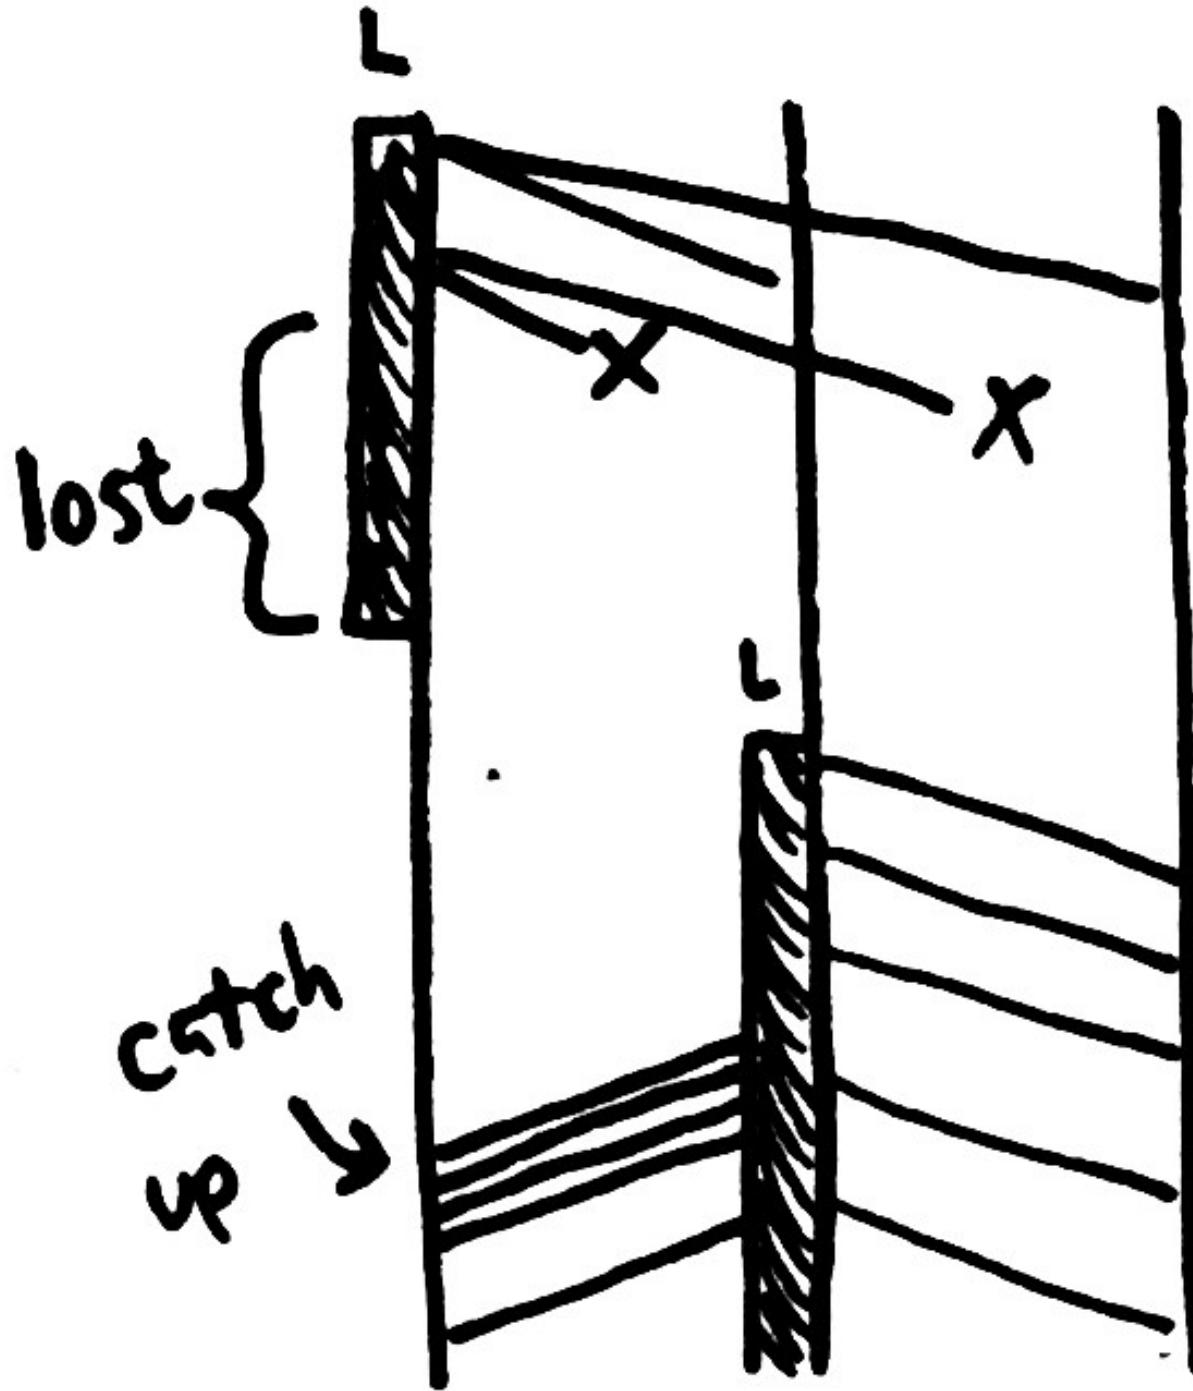


ISR

Old leader's commits
lost!



ISR



```
124 :error 10080 [██████]
Failed to send messages after 3 tries.
kafka.common.FailedToSendMessageException: Failed to send messages after 3 tries.
at kafka.producer.async.DefaultEventHandler.handle (DefaultEventHandler.scala:90)
kafka.producer.Producer.send (Producer.scala:74)
kafka.javaapi.producer.Producer.send (Producer.scala:32)
clj_kafka.producer$send_message.invoke (producer.clj:19)
jepsen.kafka$app$reify_3240.add (kafka.clj:64)
jepsen.set_app$eval2525$fn_2548$G_2515_2551.invoke (set_app.clj:12)
jepsen.set_app$eval2525$fn_2548$G_2514_2555.invoke (set_app.clj:12)
clojure.lang.AFn.applyToHelper (AFn.java:163)
clojure.lang.AFn.applyTo (AFn.java:151)
clojure.core$apply.invoke (core.clj:619)
clojure.core$partial$fn_4190.doInvoke (core.clj:2396)
clojure.lang.RestFn.invoke (RestFn.java:408)
jepsen.load$wrap_catch$catcher_2494.invoke (load.clj:110)
jepsen.load$wrap_latency$measure_latency_2491.invoke (load.clj:100)
jepsen.load$wrap_record_req$record_req_2503.invoke (load.clj:129)
jepsen.console$wrap_ordered_log$logger_2372.invoke (console.clj:108)
jepsen.load$map_fixed_rate$boss_2484$fn_2486.invoke (load.clj:65)
clojure.lang.AFn.run (AFn.java:24)
java.util.concurrent.ThreadPoolExecutor.runWorker (ThreadPoolExecutor.java:1145)
java.util.concurrent.ThreadPoolExecutor$Worker.run (ThreadPoolExecutor.java:615)
java.lang.Thread.run (Thread.java:722)
```

```
125 :ok 14668 [██████]
126 :ok 14658 [██████]
127 :ok 14659 [██████]
128 :ok 14671 [██████]
129 :ok 14662 [██████]
130 :ok 14166 [██████]
131 :ok 14157 [██████]
132 :ok 14157 [██████]
133 :ok 14165 [██████]
134 :ok 14161 [██████]
135 :ok 13660 [██████]
136 :ok 13655 [██████]
137 :ok 13657 [██████]
138 :ok 13659 [██████]
139 :ok 13655 [██████]
140 :ok 13163 [██████]
141 :ok 13152 [██████]
142 :ok 13154 [██████]
143 :ok 13152 [██████]
144 :ok 13155 [██████]
145 :ok 12660 [██████]
146 :ok 12650 [██████]
147 :ok 12654 [██████]
148 :ok 12650 [██████]
149 :ok 12651 [██████]
150 :ok 12156 [██████]
151 :ok 12148 [██████]
152 :ok 12150 [██████]
153 :ok 12151 [██████]
154 :ok 12152 [██████]
155 :ok 11653 [██████]
156 :ok 11646 [██████]
157 :ok 11648 [██████]
158 :ok 11647 [██████]
159 :ok 11650 [██████]
```

445	:	ok	4
446	:	ok	4
447	:	ok	5
448	:	ok	5
449	:	ok	5
450	:	ok	4
451	:	ok	5
452	:	ok	2
453	:	ok	5
454	:	ok	3
455	:	ok	5
456	:	ok	4
457	:	ok	5
458	:	ok	5
459	:	ok	4
460	:	ok	6
461	:	ok	5
462	:	ok	6
463	:	ok	5
464	:	ok	5
465	:	ok	5
466	:	ok	3
467	:	ok	4
468	:	ok	5
469	:	ok	4
470	:	ok	5
471	:	ok	4
472	:	ok	5
473	:	ok	7
474	:	ok	6
475	:	ok	5
476	:	ok	4
477	:	ok	5
478	:	ok	6
479	:	ok	5
480	:	ok	5

```
clojure.core$partial$fn__4190.doInvoke (core.clj:2396)
clojure.lang.RestFn.invoke (RestFn.java:408)
jepsen.load$wrap_catch$catcher_2494.invoke (load.clj:110)
jepsen.load$wrap_latency$measure_latency_2491.invoke (load.clj:100)
jepsen.load$wrap_record_req$record_req_2503.invoke (load.clj:129)
jepsen.console$wrap_ordered_log$logger_2372.invoke (console.clj:108)
jepsen.load$map_fixed_rate$boss_2484$fn_2486.invoke (load.clj:65)
clojure.lang.AFn.run (AFn.java:24)
java.util.concurrent.ThreadPoolExecutor.runWorker (ThreadPoolExecutor.java:1145)
java.util.concurrent.ThreadPoolExecutor$Worker.run (ThreadPoolExecutor.java:615)
java.lang.Thread.run (Thread.java:722)
```

664	:ok	10	
665	:ok	7841	
666	:ok	7853	
667	:ok	7856	
668	:ok	7854	
669	:ok	10	
670	:ok	7330	
671	:ok	7345	
672	:ok	7351	
673	:ok	7347	
674	:ok	8	
675	:ok	6824	
676	:ok	6842	
677	:ok	6847	
678	:ok	6841	
679	:ok	10	
680	:ok	6321	
681	:ok	6337	
682	:ok	6347	
683	:ok	6338	
684	:ok	3	
685	:ok	5814	
686	:ok	5830	
687	:ok	5840	

Kafka

100% acknowledged

50% false positives

0% false negatives

2 well-timed
failures on a
single node
mean lost data.

- ISR must shrink
- All ISR nodes must lose ZK connection.
- "All" could be 1 node.

How do you recover
tolerance to more
failures? Use majority.

-ISR lower bound of $N/z + 1$

Majority quorums are
probably optimal for NDC
Consensus fault-tolerance.

— Peleg & Wool, '93

Kafka is considering
stricter bounds on
correctness; must balance
performance.

Kafka strategies

- 0.8 is still beta
- If present, turn on stronger safety rules
- Roll your own replication?

Cassandra!

Dynamo system

- Partitioned hash ring
- Quorums & handoff
- LWW only, not vclocks

"But Last Write Wins
is bad, right?"



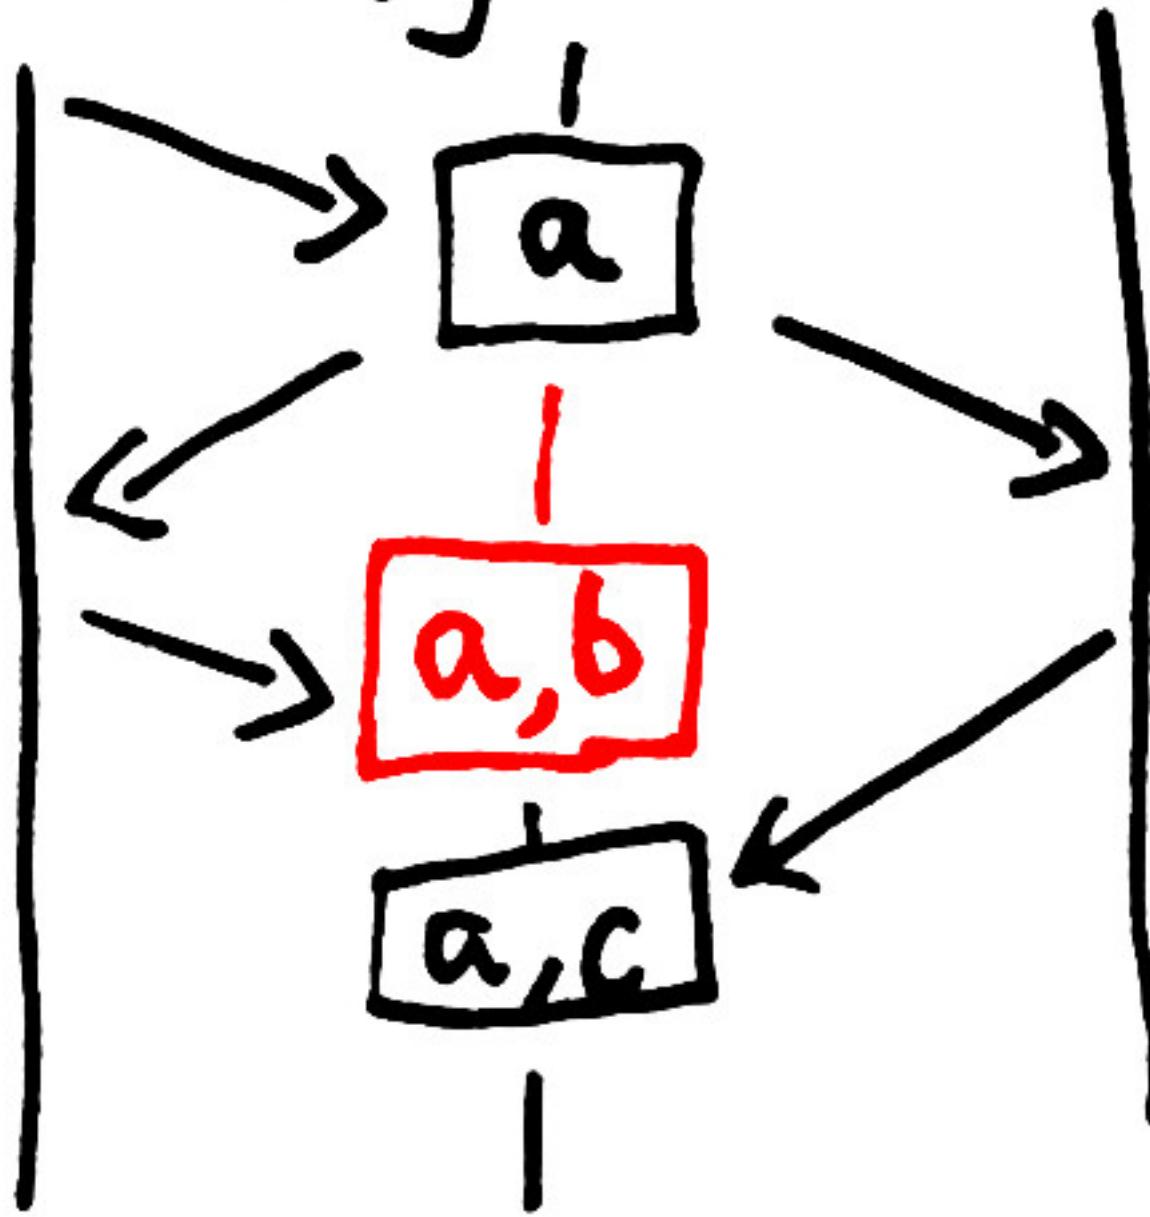
Well Sally, it depends!

You can't safely change
an LWW register.

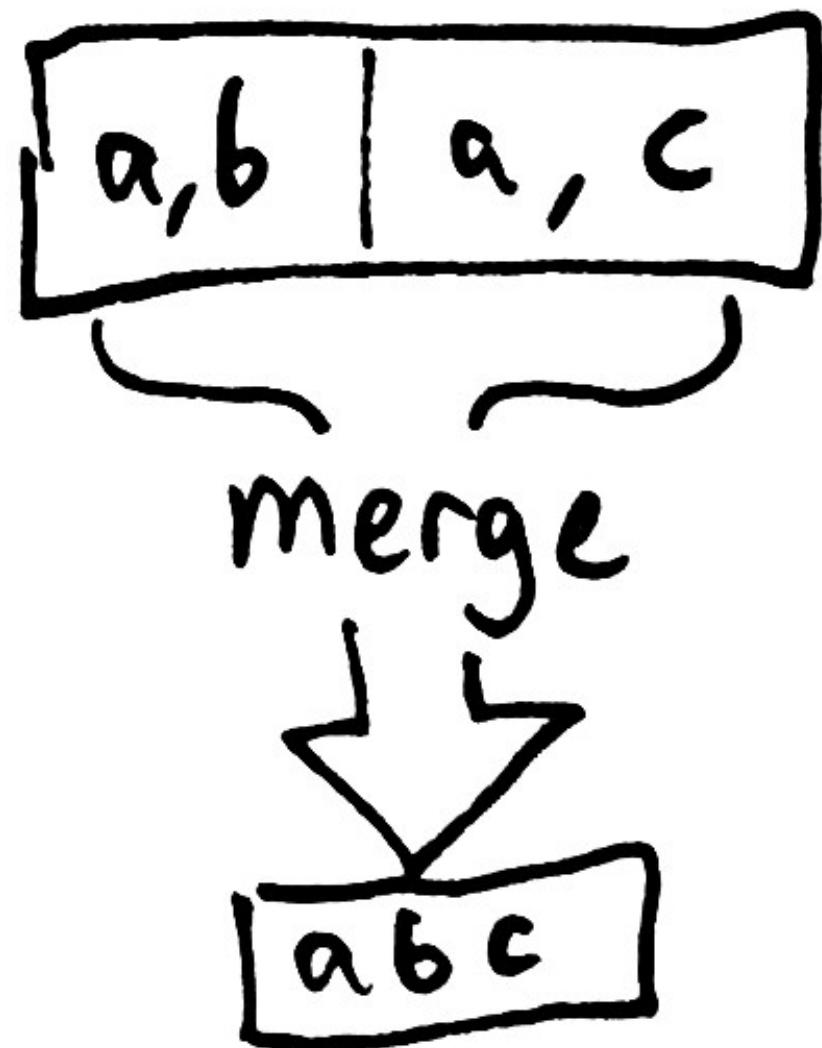
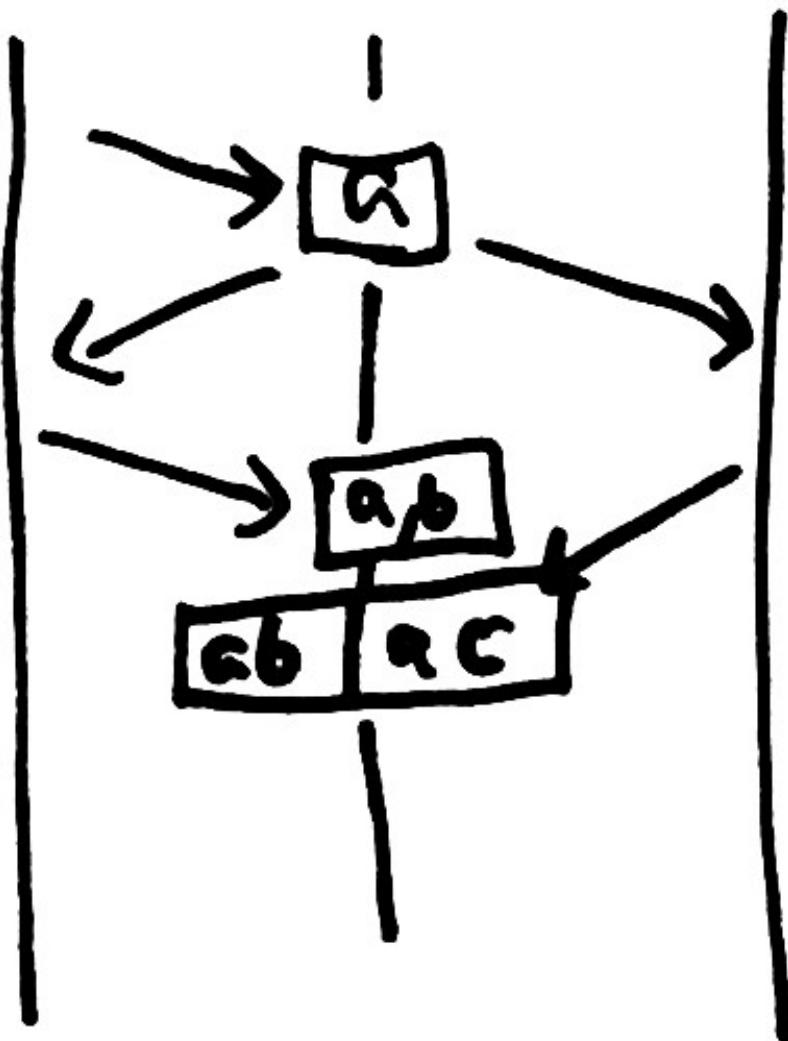
Safe?

- Guarantee that
your write is
causally connected
to a future read.

register



Vector clocks identify
these causally isolated
writes, and give you
both copies.



Merge fn must be

- Associative
- Commutative
- Idempotent

CRDTs!

Early in Cassandra history, chose not to use vclocks for performance.

Consequently, no safe way to modify a cell.

Cassandra Cell Change (quorum r+w, locks)

44% acknowledged

27% false positives

0% false negatives

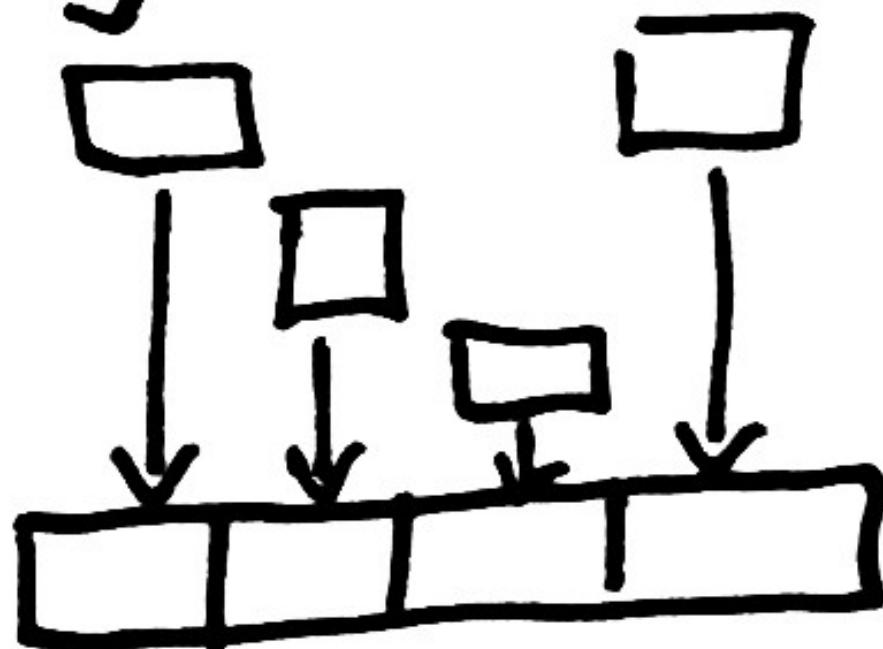
or tables. Cassandra supports tuning availability and consistency, and always gives you partition tolerance. Cassandra can be tuned to give you strong consistency in the CAP sense where data is made consistent across all the nodes in a distributed database cluster. A user

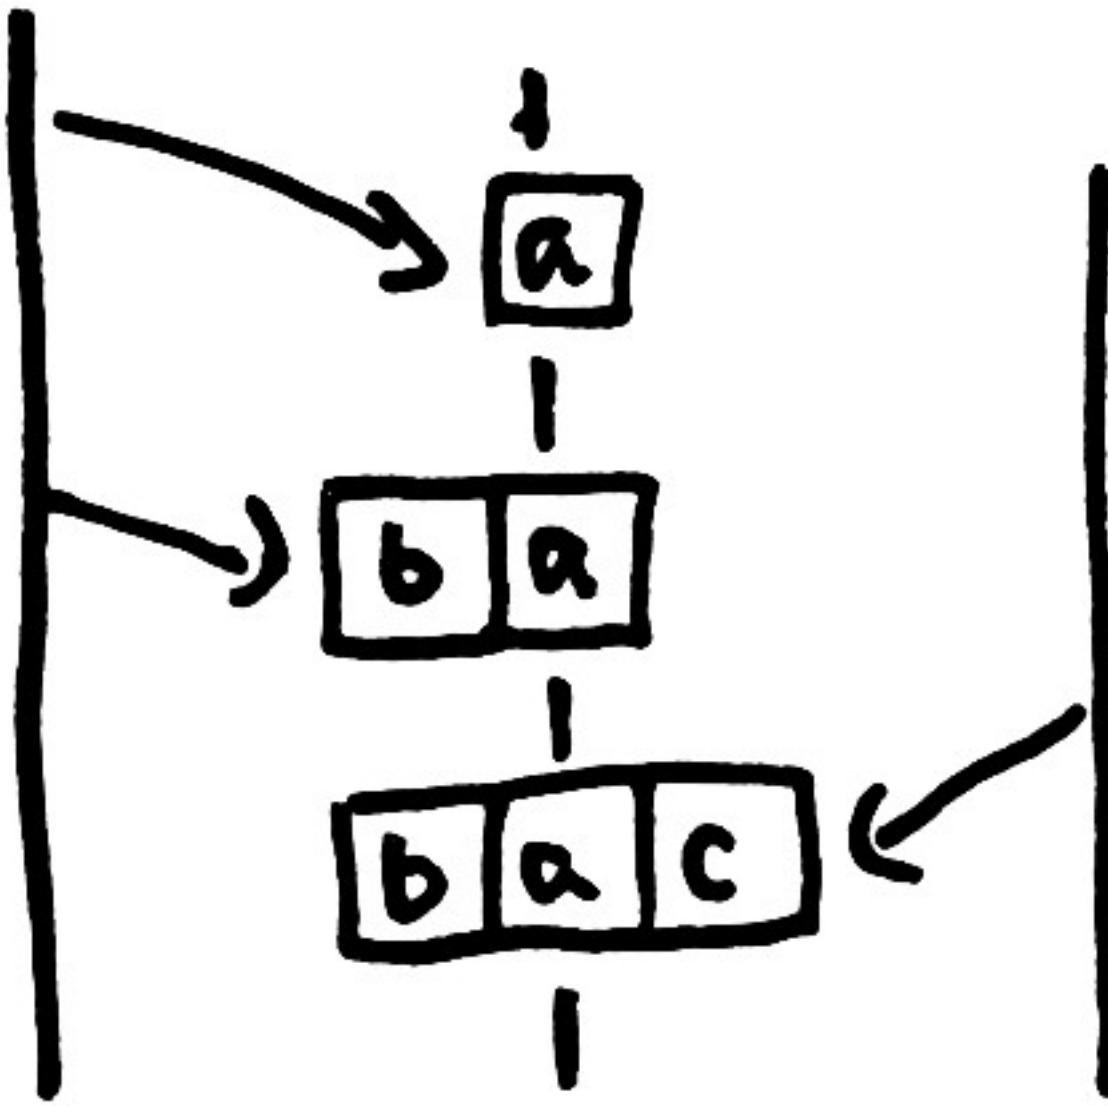
Instead, Cass. has evolved efficient ways to write every change to a distinct cell.

Immutable operation

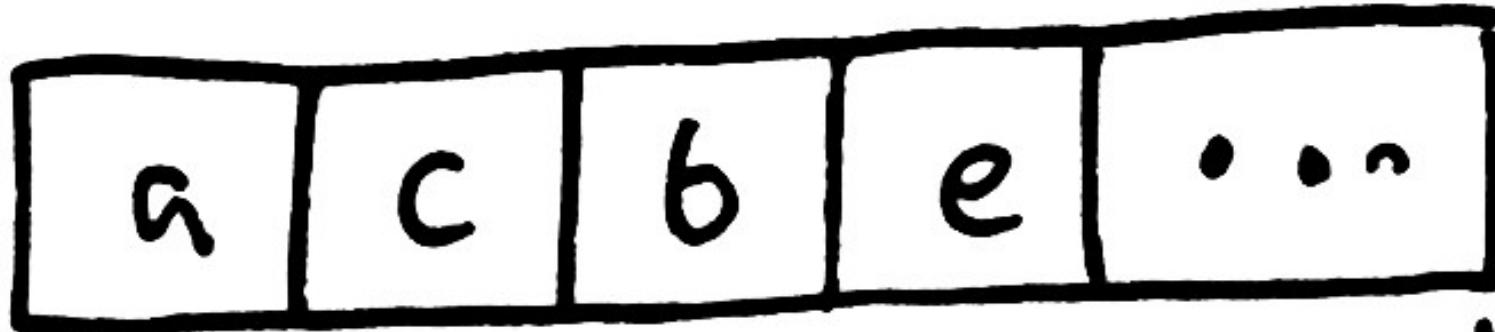
log!

row :





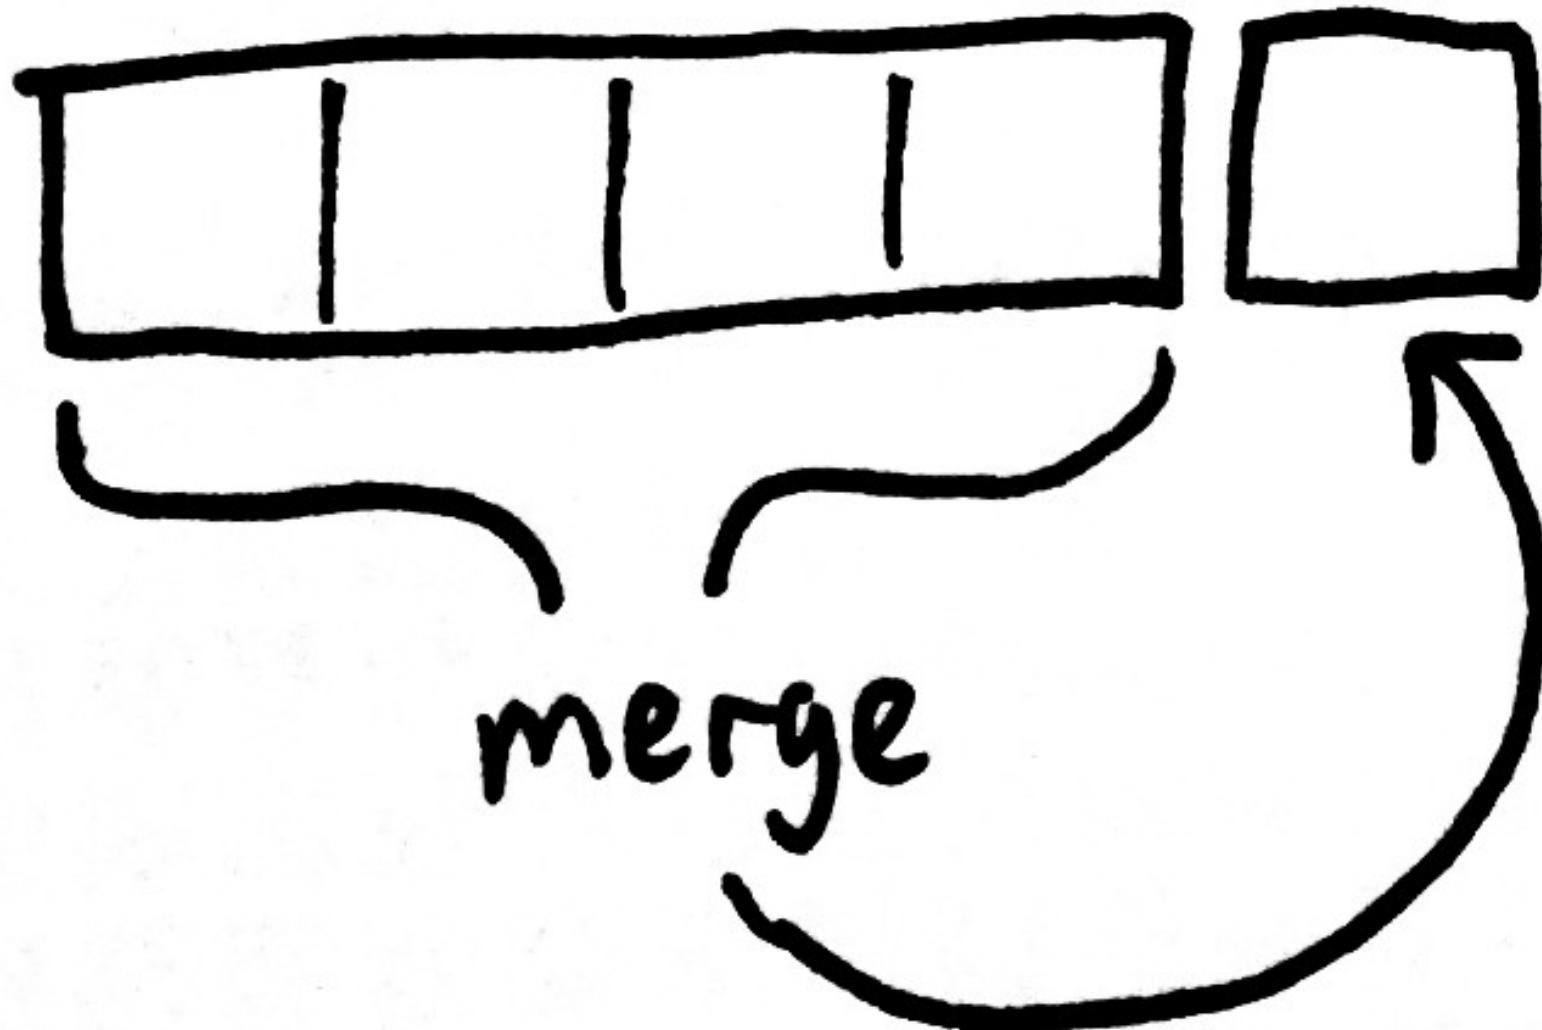
Reads require
merge
functions



merge



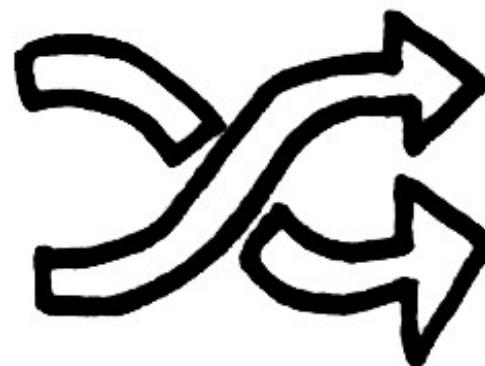
Garbage Collection?



~~Obsolete~~

Vector clocks handle
identifying old, unneeded
data.

Since there is no
global ordering for
updates, still need
commutative merges.



So! Cassandra
and Riak have
similar consistency
semantics. CQL
provides limited
CRDTs.

```
85 :ok 4
86 :ok 4
87 :ok 4
88 :ok 6
89 :ok 6
90 :ok 6
91 :ok 6
92 :ok 7
93 :ok 8
94 :ok 5
95 :ok 5
96 :ok 5
97 :ok 5
98 :ok 12
99 :ok 12
```

Initiating partition.

```
100 :ok 2
101 :ok 6
102 :ok 6
103 :ok 6
104 :ok 6
105 :ok 4
106 :ok 3
107 :ok 5
108 :ok 7
109 :ok 3
```

Partitioned.

```
110 :ok 4
111 :ok 6
112 :ok 6
113 :ok 5
114 :ok 7
115 :ok 4
116 :ok 3
117 :ok 4
118 :ok 4
119 :ok 4
120 :ok 6
121 :ok 8
122 :ok 7
123 :ok 7
124 :ok 8
125 :ok 7
126 :ok 7
127 :ok 10
128 :ok 9
129 :ok 11
130 :ok 6
131 :ok 6
132 :ok 6
133 :ok 7
```

Cassandra CQL sets

100% acknowledged

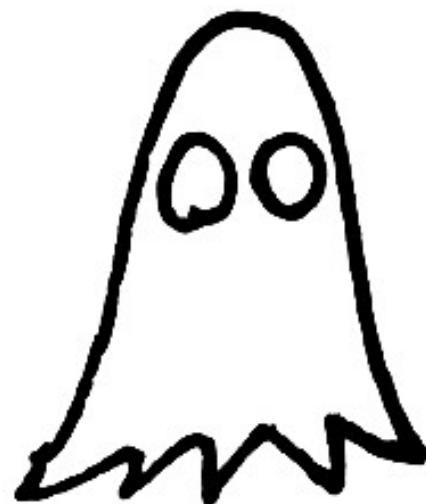
0% false positives

0% false negatives

CQL collections
are partition tolerant!

- But semantics are
subtle...

BEWARE DELETES



Cassandra counters

at consistency = quorum

+/- 50% actual value

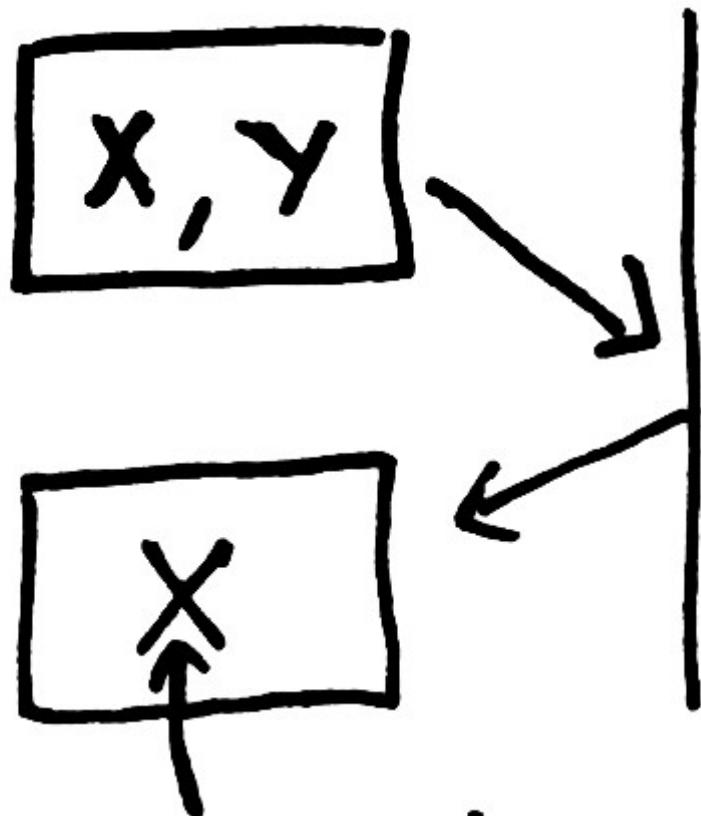
What about
ISOLATION?



"From a transactional ACID...
standpoint, [atomic row ops]
give Cassandra transactions
ACID support"

- Datastax "About Writes"

"Atomic"

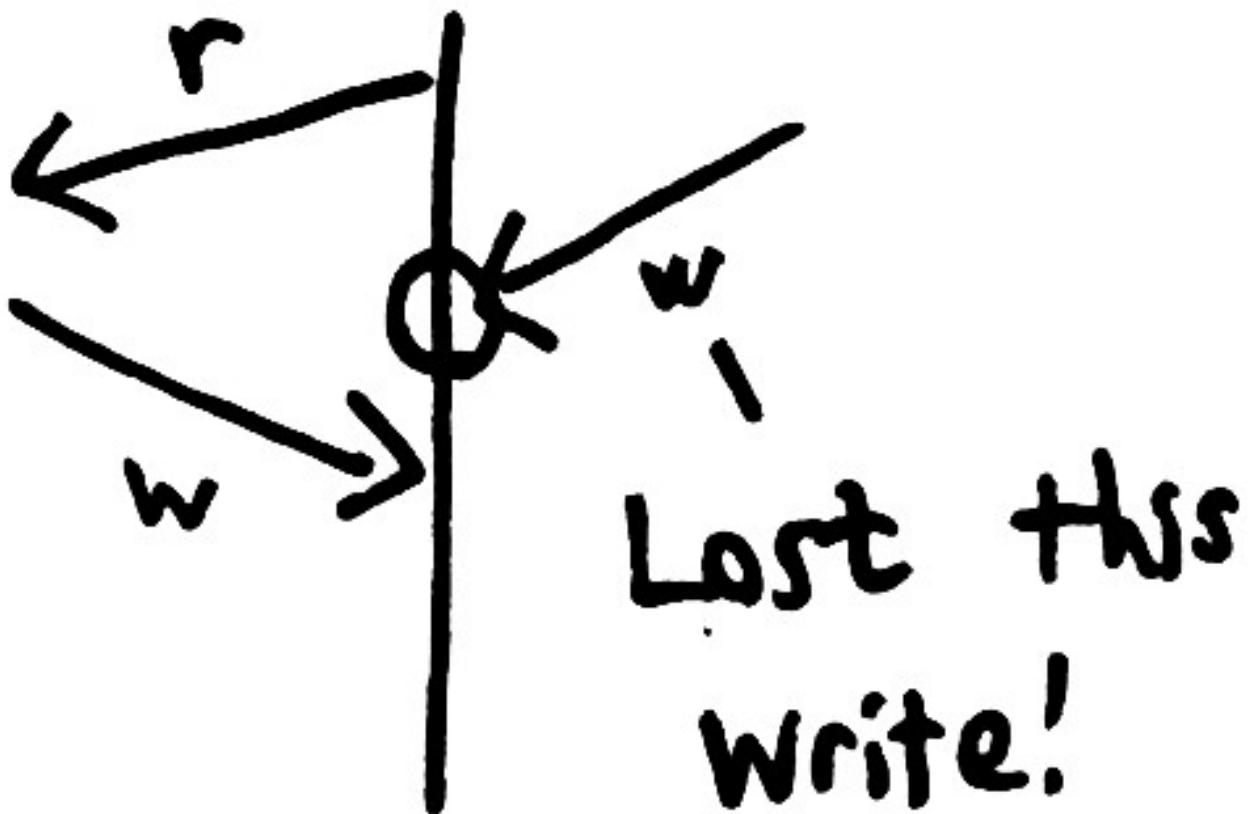


Wrong!

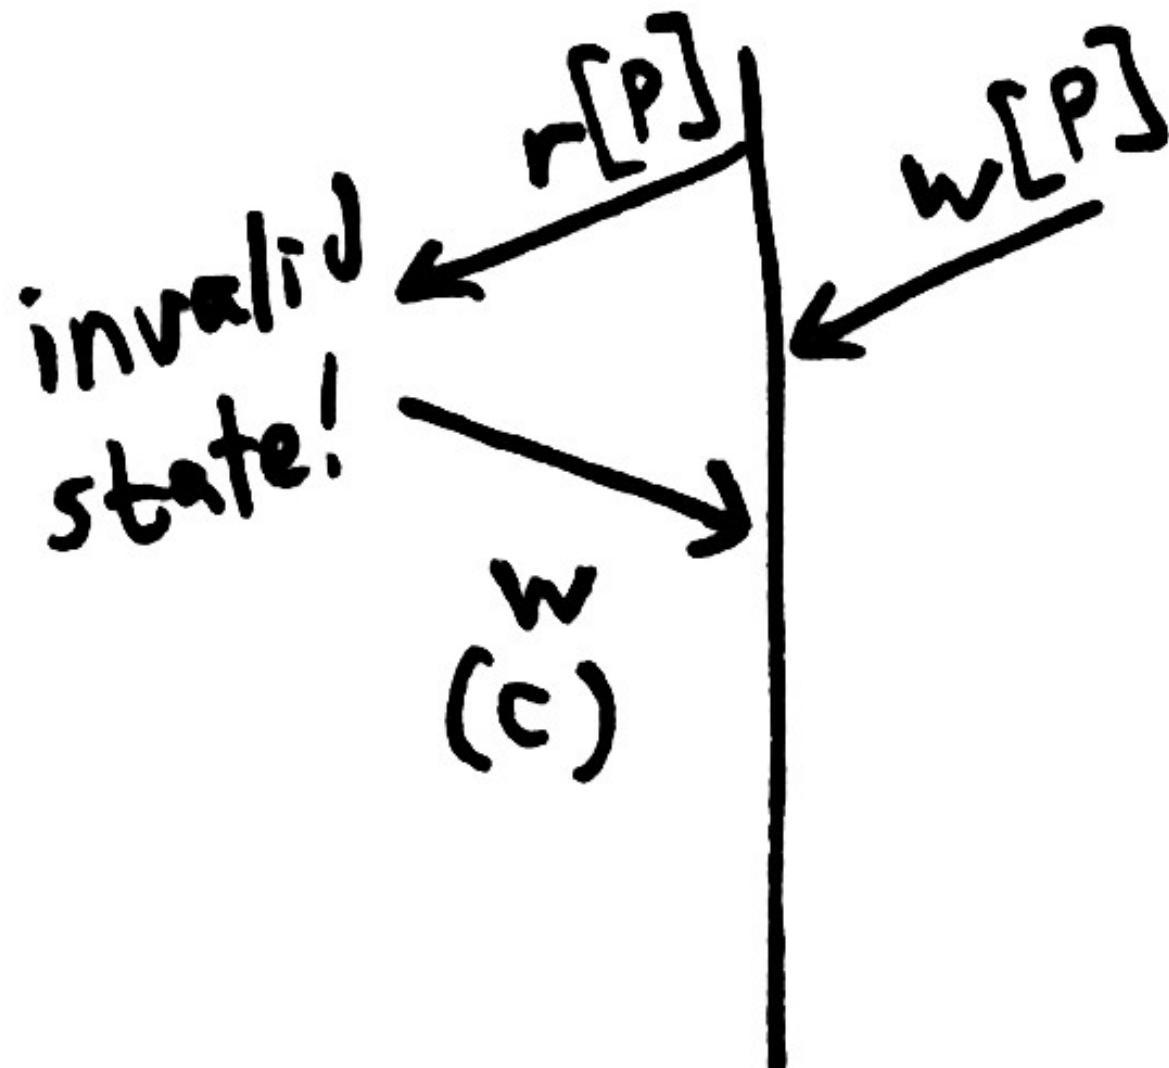
"Isolated"

... um ...

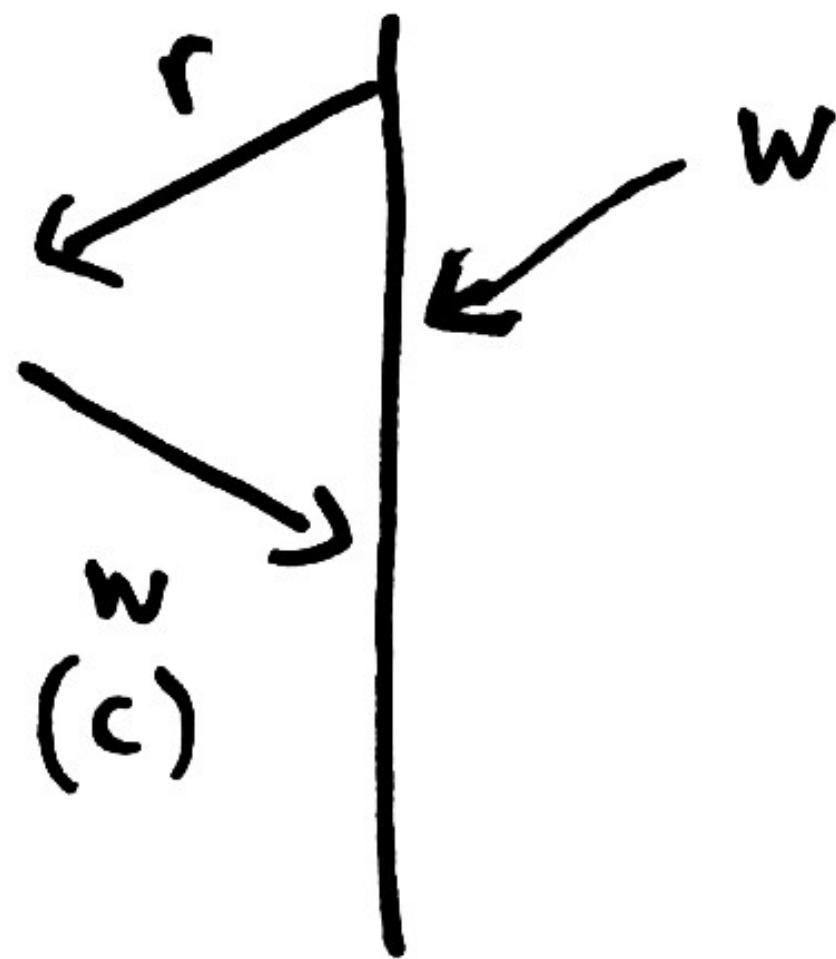
P4: Lost update (the LWW problem!)



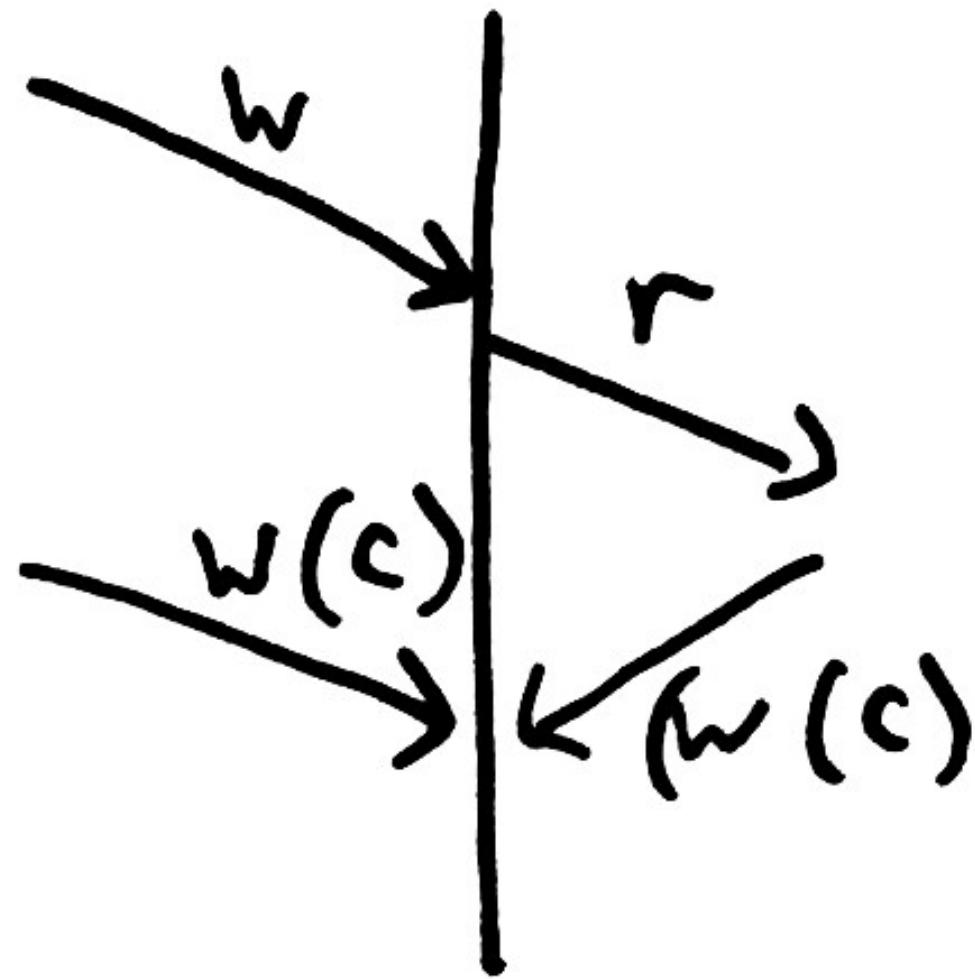
P3; Phantom



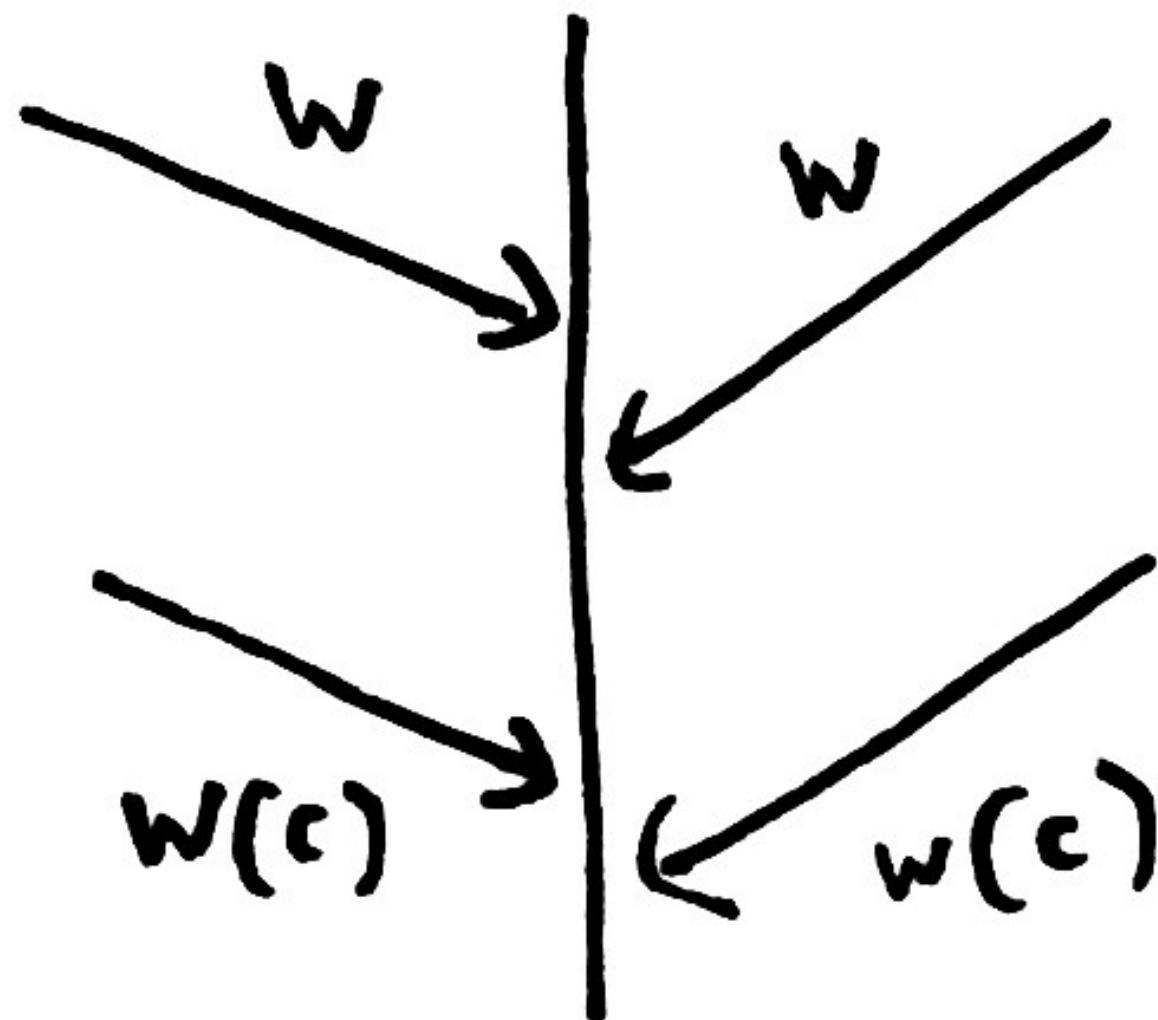
P2: Fuzzy Read



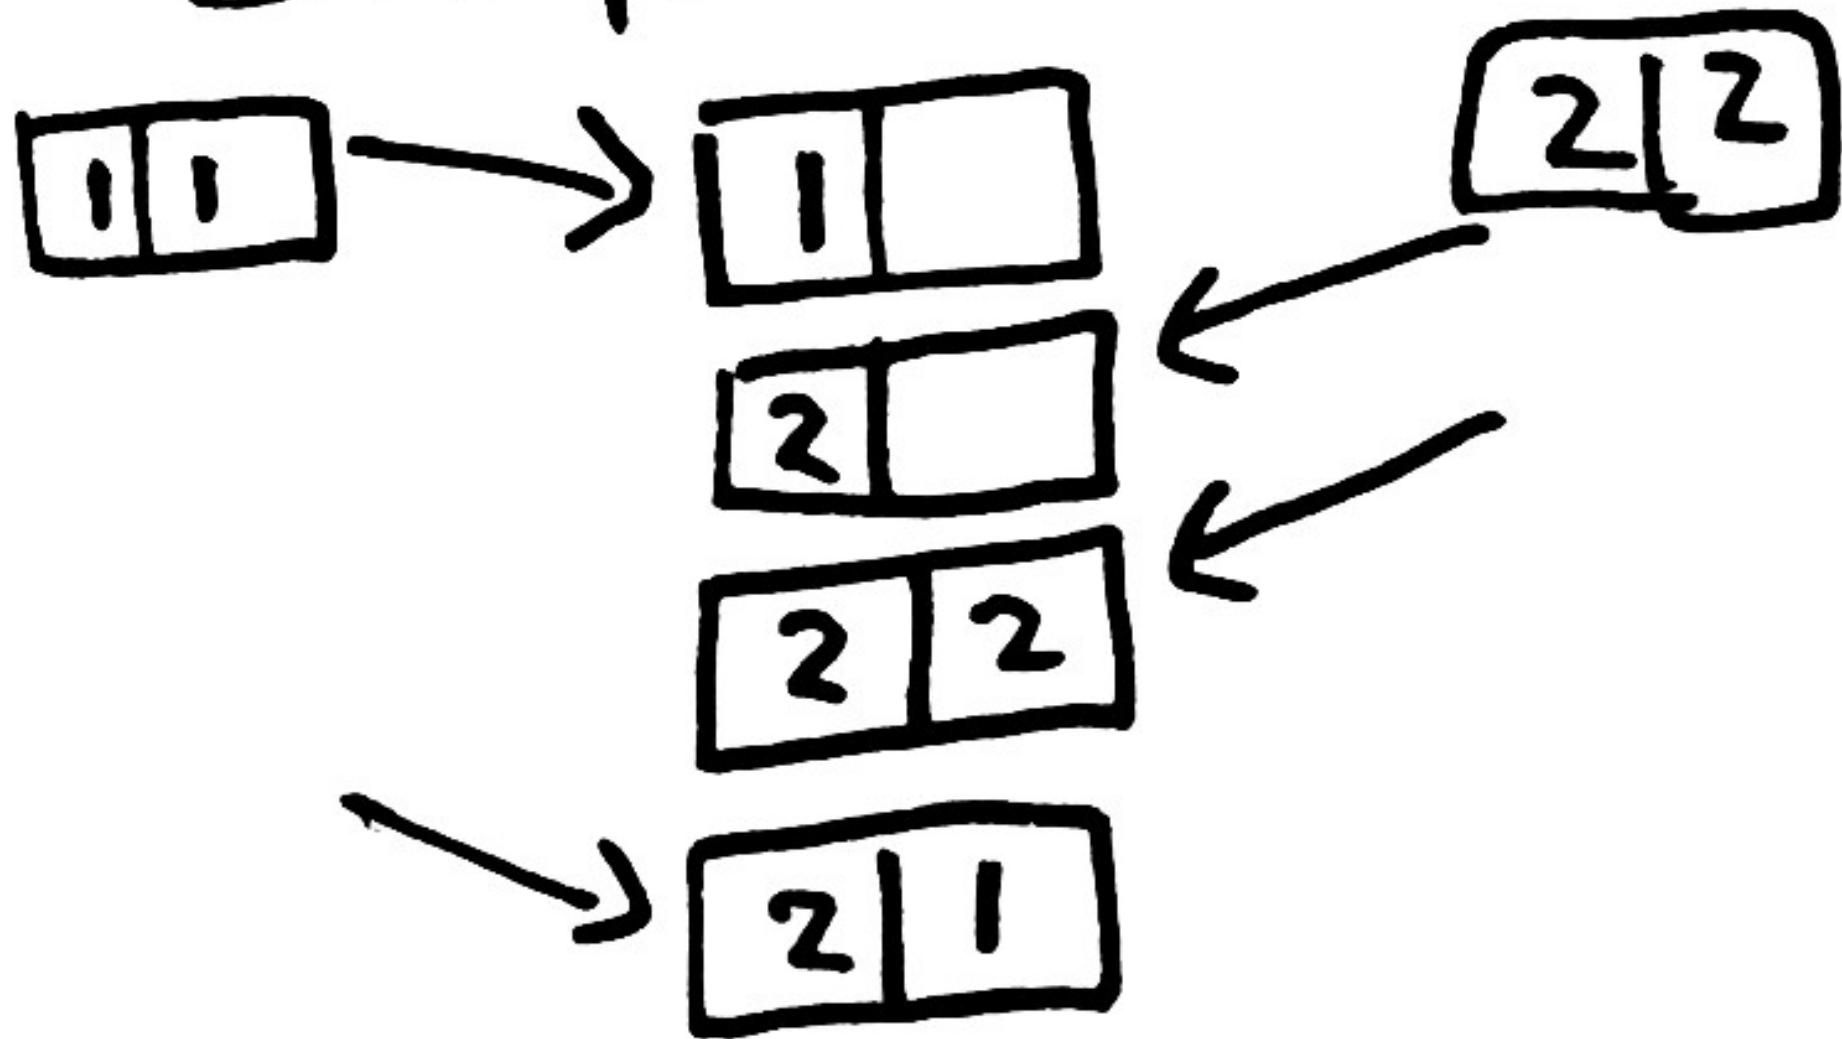
P1: Dirty Read



PQ: Dirty write



Example of P0:



"ANSI SQL Isolation
should ... require PQ
for all isolation levels"

-Berenson, et al

MSR -TR -95 -51

Cassandra 1.1 as row-level updates are now made in isolation. Cassandra 1.1 guarantees that if you update both the login and password in the same update (for the same row key) then no concurrent read may see only a partial update.

From a transactional ACID (atomic, consistent, isolated, durable) standpoint, this enhancement now gives Cassandra transactional AID support. A write is isolated at the row-level in the storage engine.

Actually, they are isolated in the sense of ACID,

Certainly there is nothing like ACID, but a single multi-cell update to a given record is guaranteed to be atomic, and if you have two concurrent multi-cell updates to a single record, they are guaranteed to eventually resolve to a consistent ordering of those operations

The difference is that all the cells with the same 'id' would belong to the same partition and stored together, so you'd be able to write them atomically and read them together cheaply in a single operation.

It wasn't really mentioned in the article but if you do both {X1,Y1} and {X2,Y2} as full row updates then no you can't end up with a mixed row.

<http://www.datastax.com/dev/blog/row-level-isolation>

In Cassandra the second update agreed upon by the cluster would "win", so either (X1,Y1) or (X2,Y2).

Cassandra

allows

POLY

(and more!)

In Cassandra,
write order doesn't
matter. LWW
always takes prio.

What if timestamps
are equal?

$$t_a = t_b$$

Compare values

$a < b$

winner

1	-1
---	----

2	-2
---	----

-	-2	-1
2		

2	-1
---	----

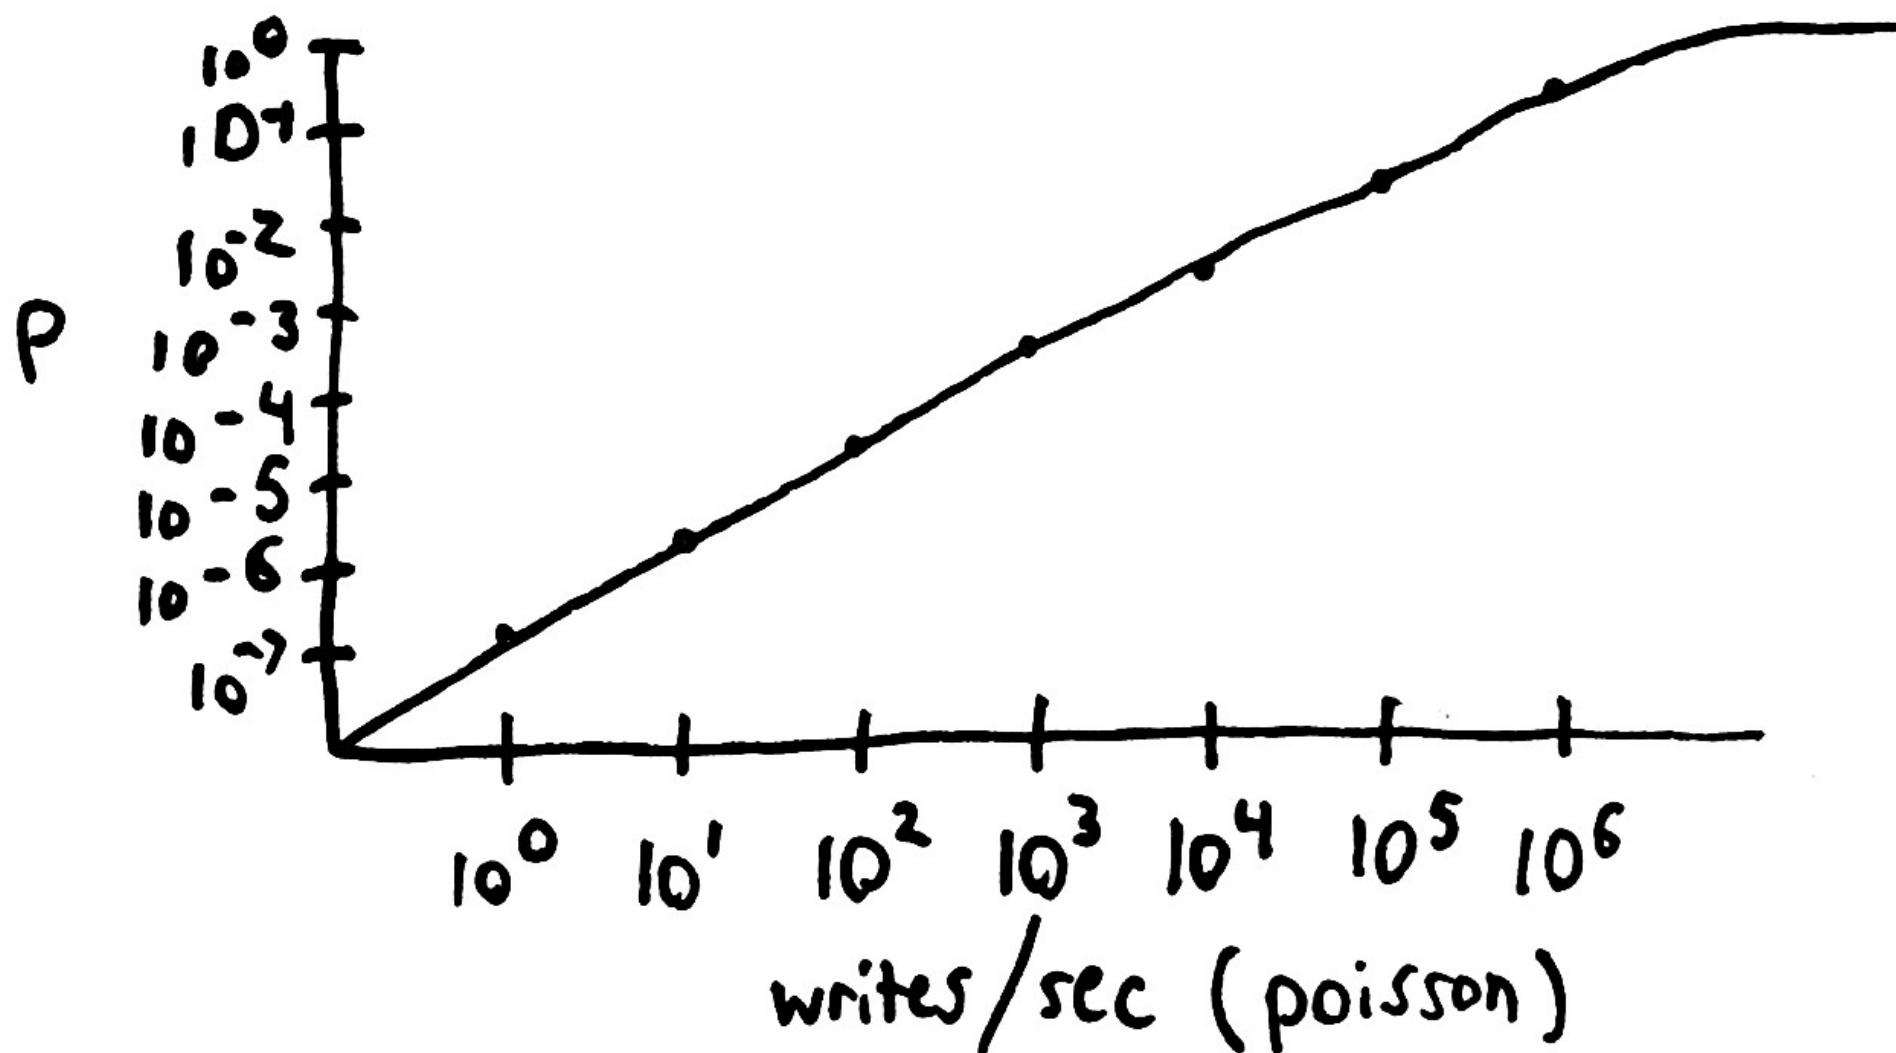
"But timestamps
never conflict!"

"But timestamps
never conflict!"

It Depends!TM



Probability of a microsecond
conflict visible in a read:



10 writes/sec

1 read/sec

34% chance of
collision in
a day

2 writes per row,
separated by mean

latency 100 ms

(poisson processes)

$$P(\text{corrupt}) = 5 \times 10^{-6}$$

Theoretical
Limit!

With Cassandra

1.2.8 & 2.0.0,

Jars driver 1.0.3,

2 writes / row

mean latency 100ms

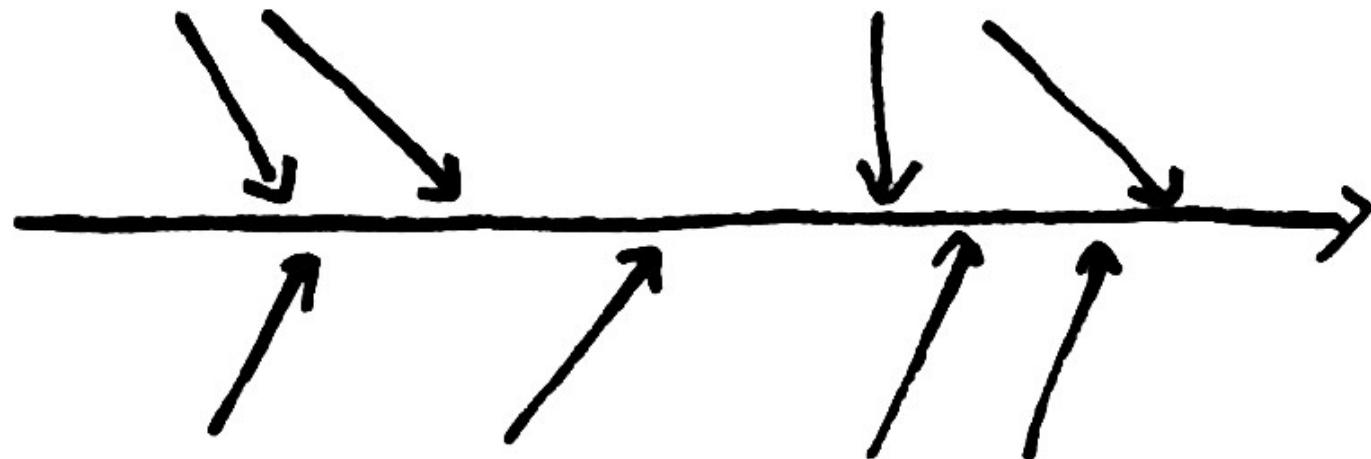
1 in 250

rows found

Corrupt!

USE
lightweight
TRANSACTIONS

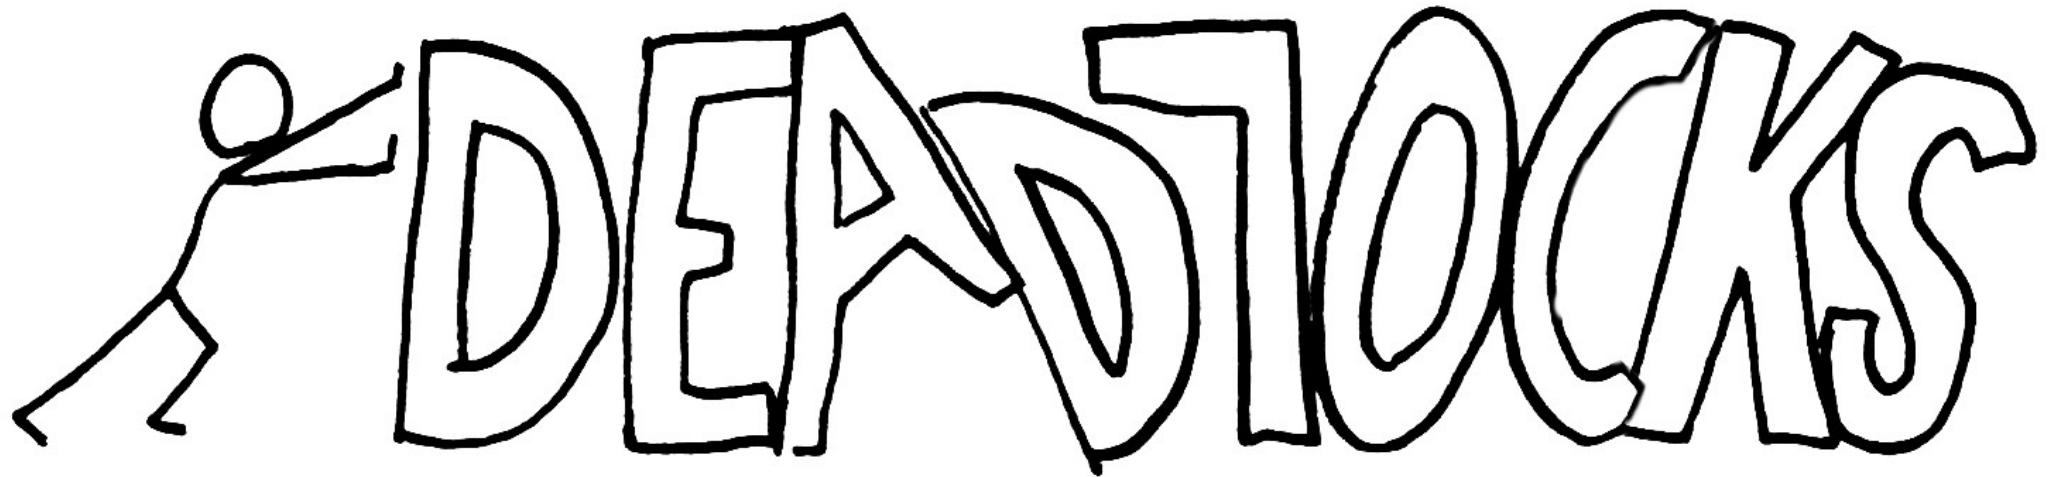
"Linearizable Consistency"



- DataStax, "Lightweight
Transactions in Cassandra

20.0'

Java driver
does not
support PRXas
transactions.

A black and white line drawing of a stick figure pushing a large, stylized word. The word is composed of thick, outlined letters. The first letter, 'D', is partially cut off on the left by the edge of the frame. A stick figure is positioned on the far left, leaning forward and pushing against the vertical stroke of the 'D'. The word itself is 'DEADLOCKS'.

Paxos txns lock
up the cluster
permanently.

-clear system.paxos

\$ git checkout paxos-fixed-hopefully

Paxos transactions
are not linearizable.

#6012

#6013

Can accept two
values for same
paxos round.

Failed proposals can
still succeed in later
rounds.

Cassandra Transactions

5% acknowledged

4% false positives

9% false negatives

"Paxos made without
stress tests."

Open source has had the reputation of producing good imitations, but not innovation. Perhaps Cassandra's origins as a hybrid of Dynamo and Bigtable did not disprove this, but Apache Cassandra's development of lightweight transactions and CQL are true industry firsts.

Cassandra strategies

- Use as a strict AP store
- No isolation/atomicity guarantees for writes
- Avoid transactions

to

Re CAP

Jepsen only
tells you about
one system

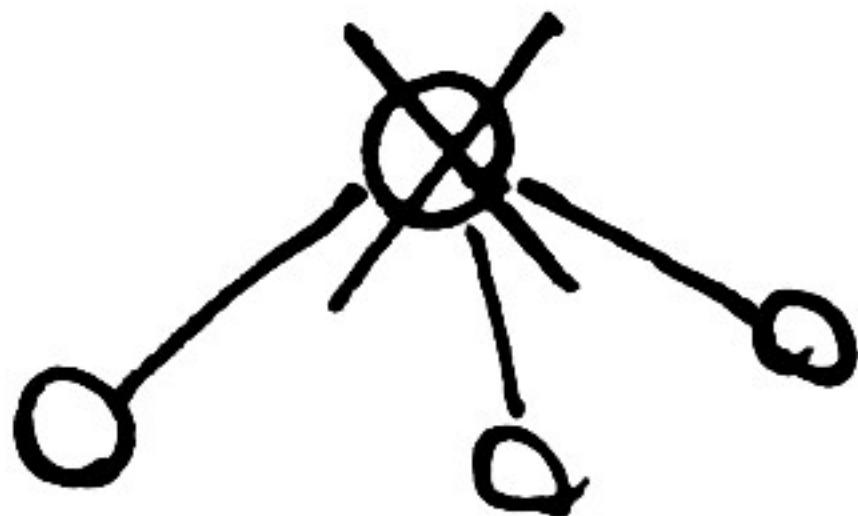
Clients are
a part of the
distributed system!

- Clients
- Servers
- Network
- Semantics

Measure
your
Systems

ASK:

"What if we
lose this node?"



"What if these
nodes pause for
3 minutes?"

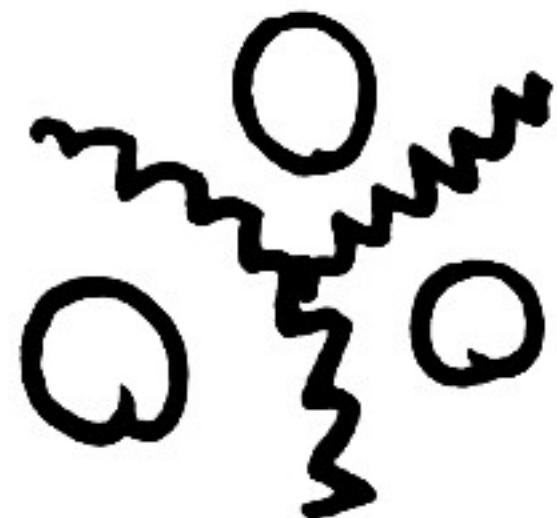


"How do we rely
on clocks?"

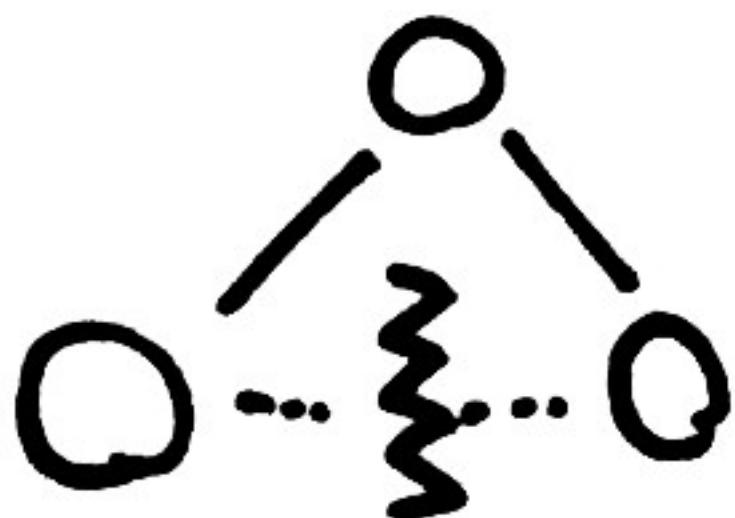


"What if the
network delivers
this message 5
minutes late?"

"What if there
is no majority?"



"What about partial
visibility?"



"What safety and
liveness guarantees
does this system
really need?"

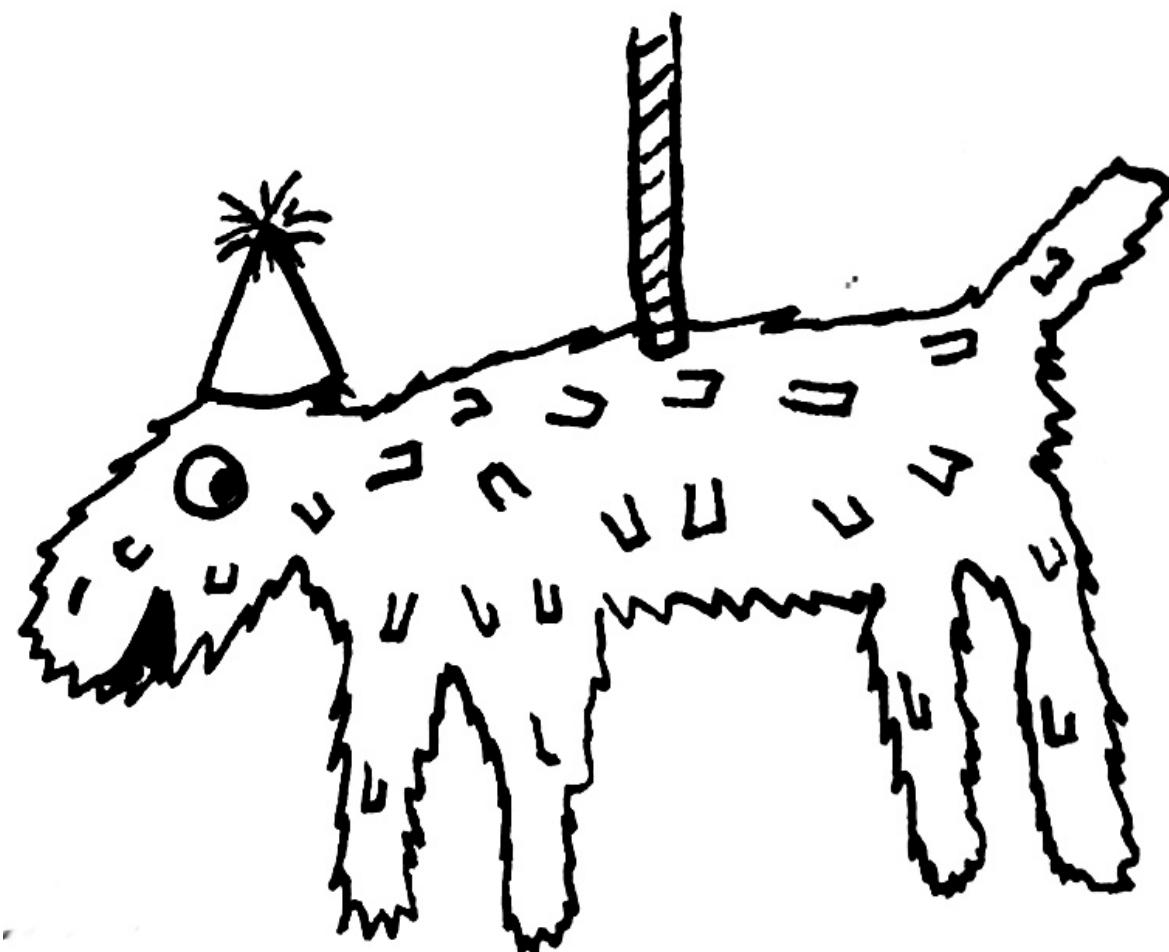
Data loss or
inconsistency can
be OK!

Balance

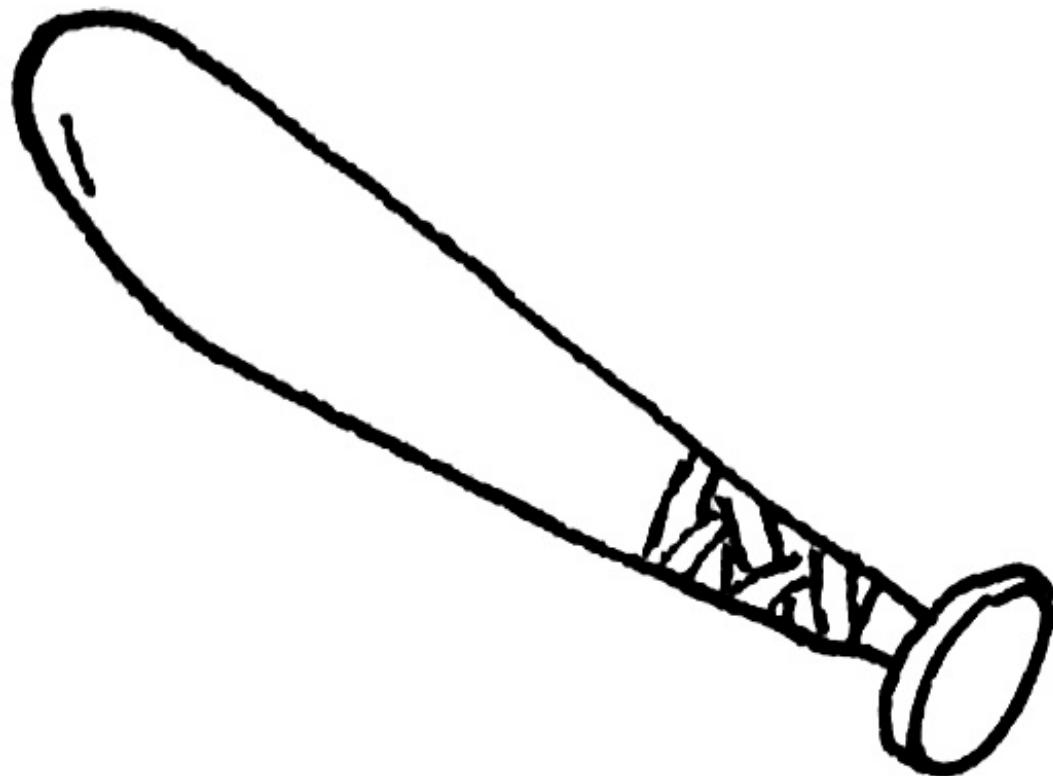
Complexity

— with —

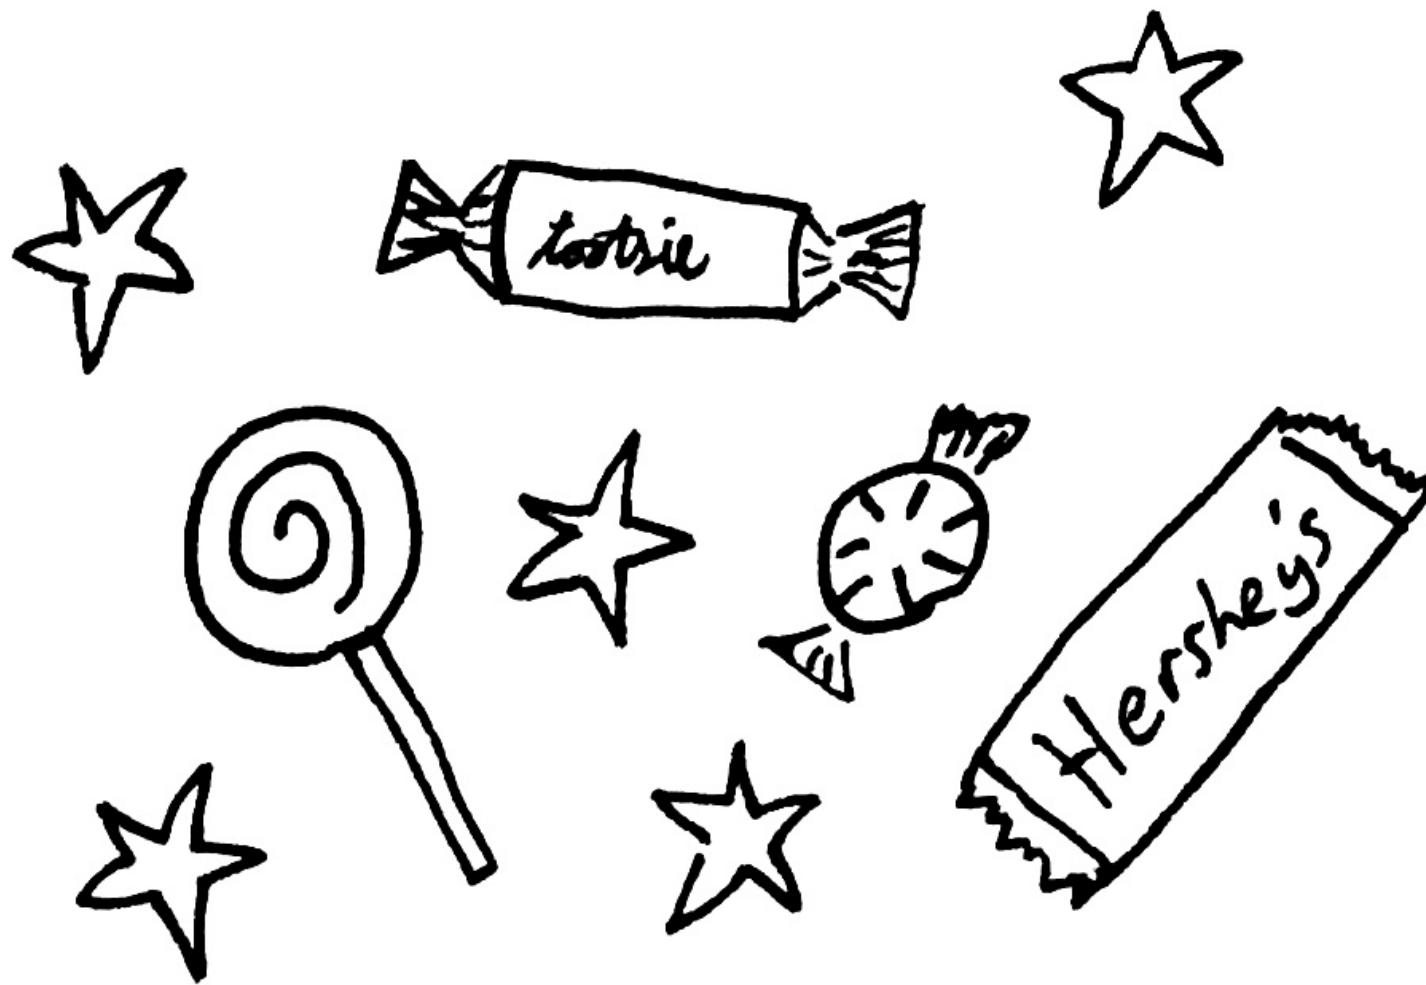
Correctness



Write a
test program!



Cause
Failures



SEE WHAT
HAPPENS!

Thank You!

Jonathan Ellis

Aleksey Yeschenko

Drew Blas

Rick Branson

Seth Proctor

Michael Klishin

Eric Lindvall

Duke Adamonis

Camille Fournier

leftside

Neha Narkhede

Jay Kreps

Patrick McFadin

Jun Rao

Joseph Blomstedt

Kelly Sommers

Peter Baillis

Thanks!

github.com/aphyr/jepsen